



High Impact Skills Development Program

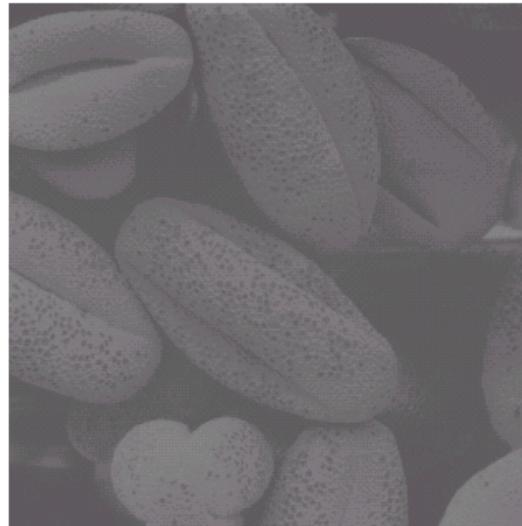
in Artificial Intelligence, Data Science, and Blockchain

Module: [Computer Vision]
Lecture 2: [Image Preprocessing]

Instructor: [Dr. Rabia Irfan]
[Assistant Professor], SEECS, NUST



Image Enhancements



Process of improving the visual quality of an image by modifying its features to make it more suitable for specific applications or more visually appealing for human observation.

Two types:

→ Spatial

(**Geometric** + Intensity) –

Individual Pixel based

→ Spatial Filtering – Pixel Neighborhood based

Transformation



Agenda



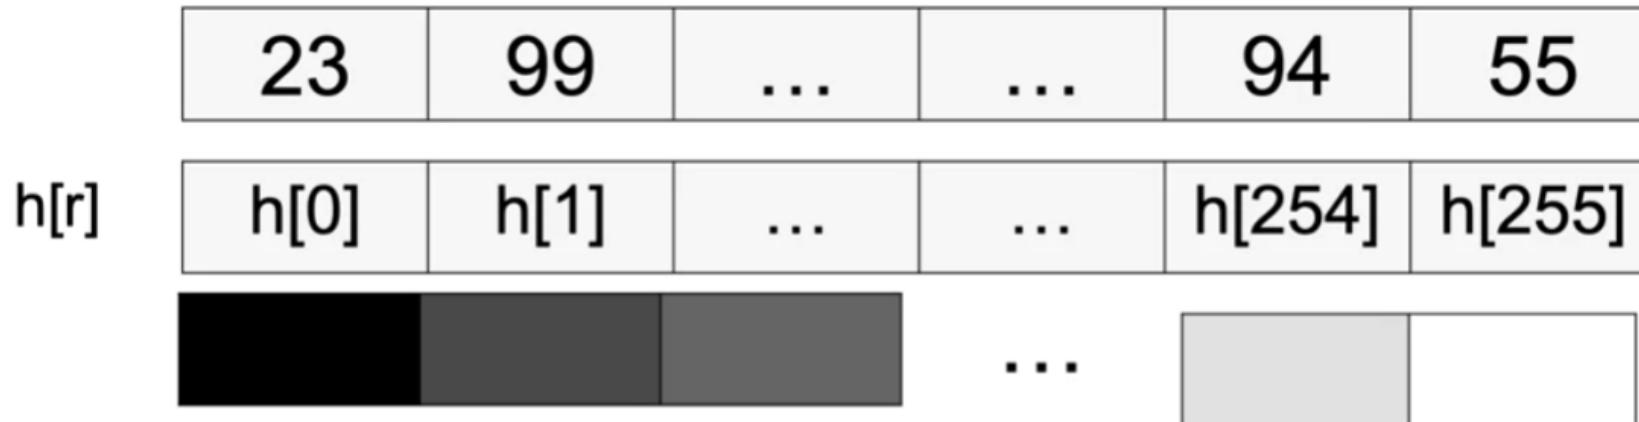
- Following topics will be discussed under the topic of Image Enhancement:
 - Histograms
 - Intensity Transforms
 - Spatial Filtering



Histograms



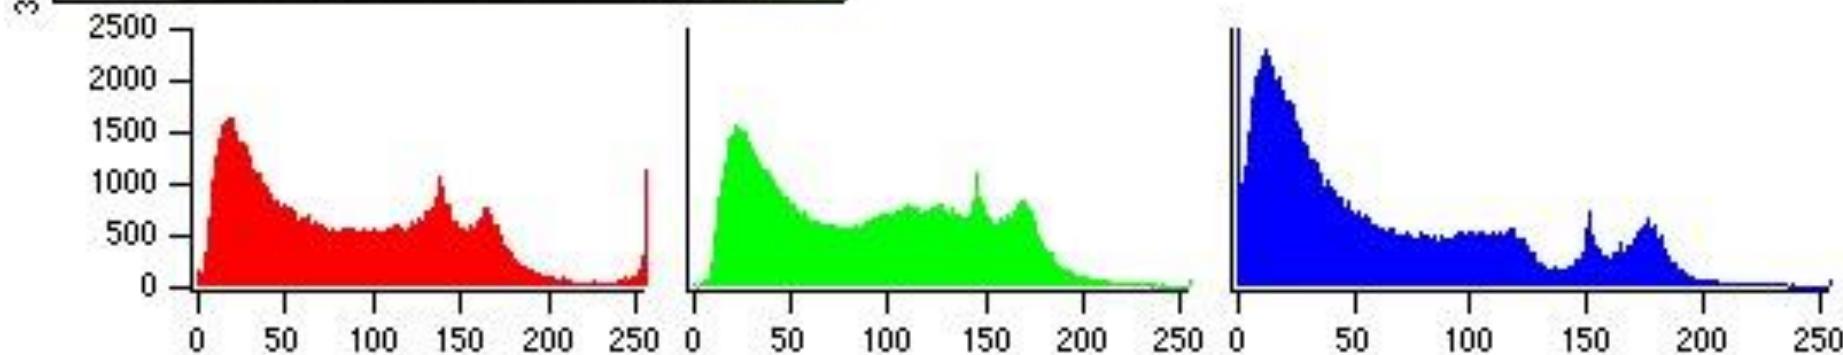
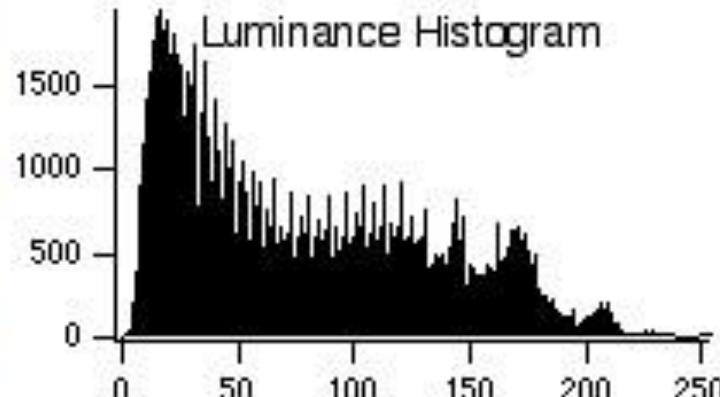
- Counts the number of occurrences or frequencies of pixel intensities in an image



- It plots the number of pixels vs. each intensity value.

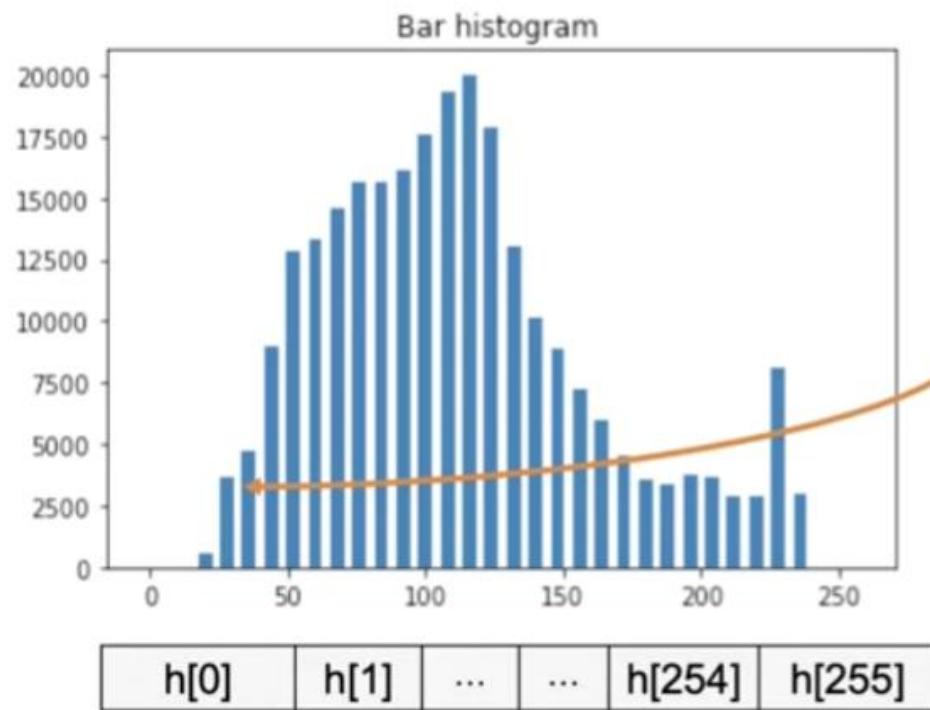


Histograms





Histograms





Histograms



```
import cv2
```

```
goldhill = cv2.imread("goldhill.bmp")
```

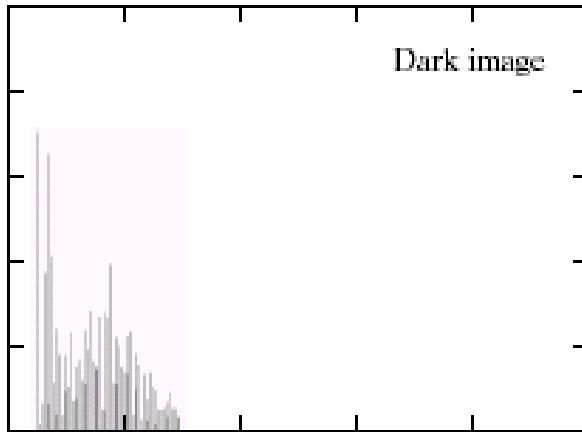
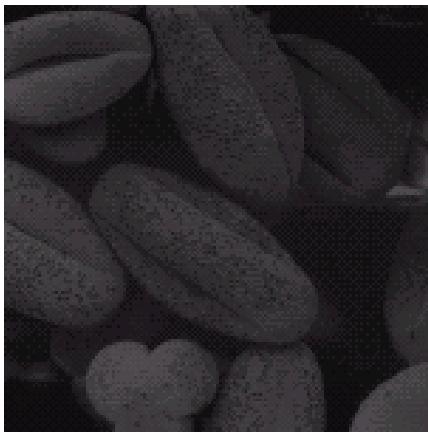
```
hist = cv2.calcHist([goldhill],[0],None,[256],[0,255])
```

34	1	94	0
0	1	2	...	254	255

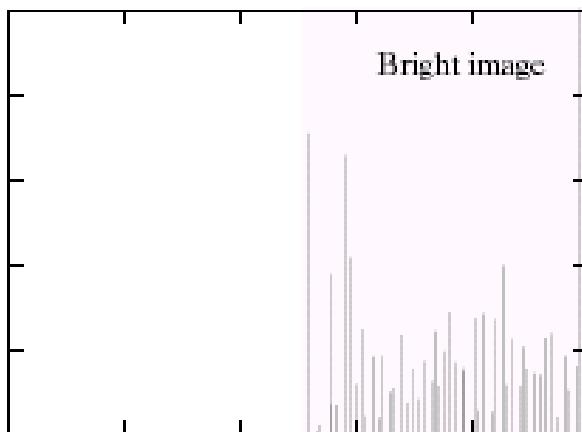
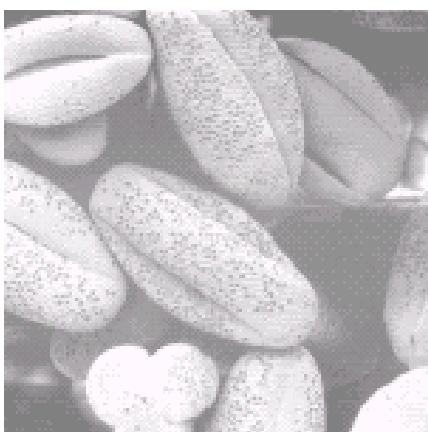




Histograms



Dark Image

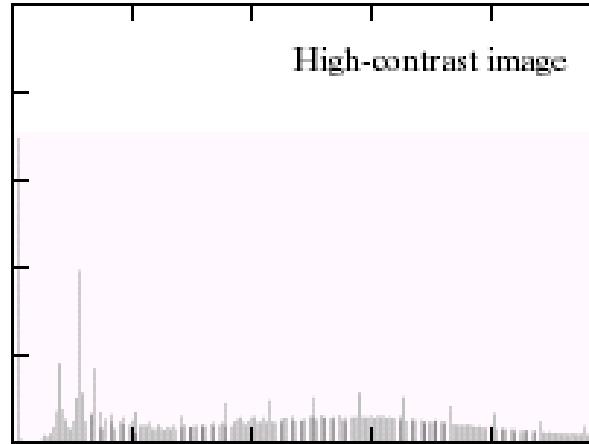
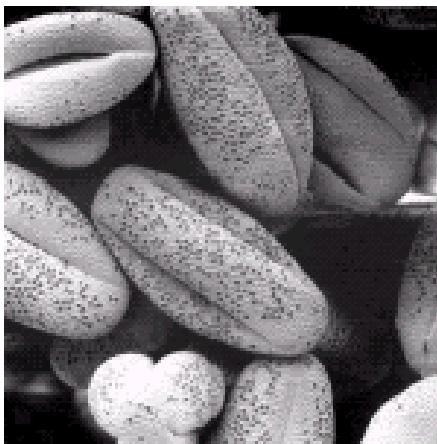
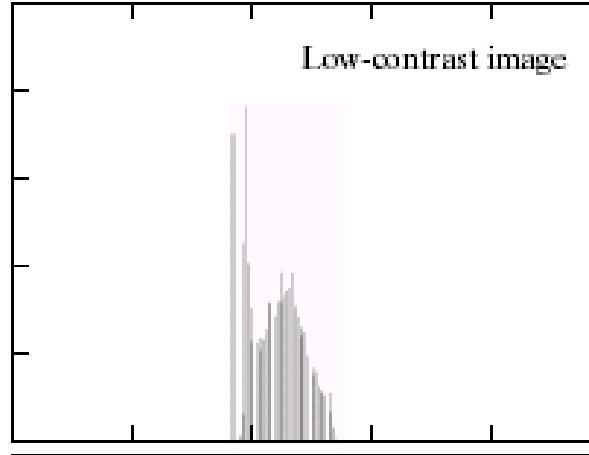
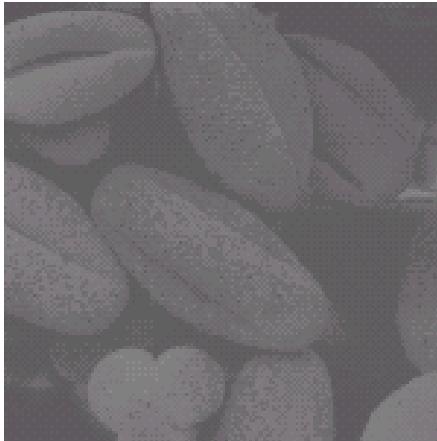


Bright Image

How would the histograms of these images look like?



Histograms



Low Contrast Image



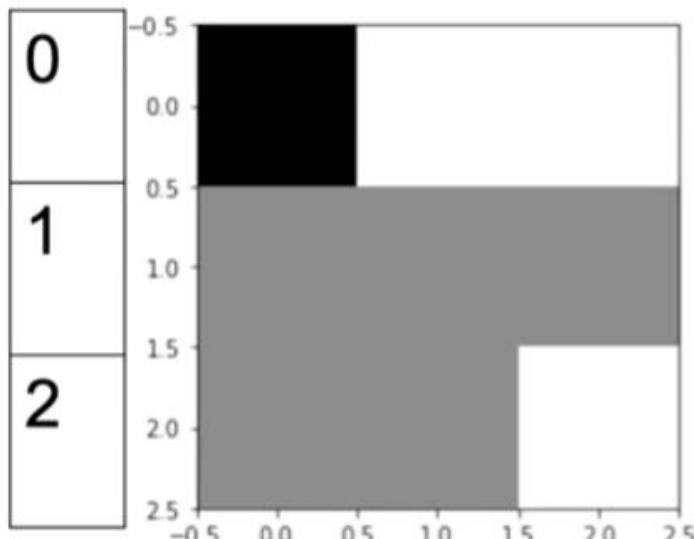
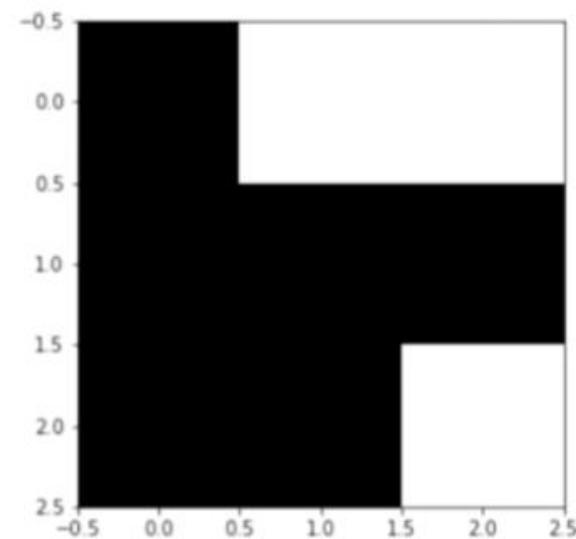
How would the histograms of these images look like?

High Contrast Image





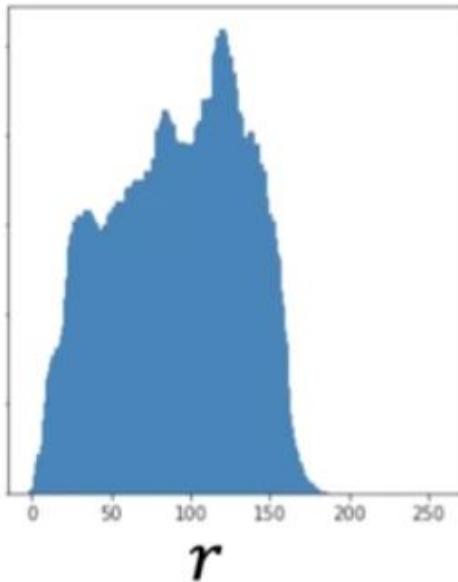
Intensity Transformations

 $f[i, j]$  $g[i, j] = T\{f[i, j]\}$  $g[i, j]$  i, j

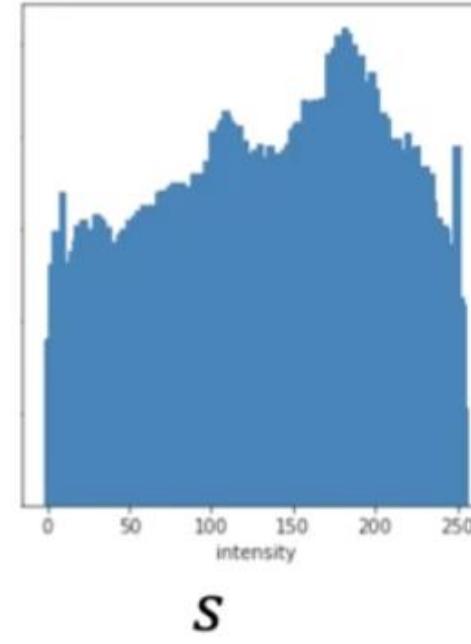
0	1	2
---	---	---



Intensity Transformations


$$h_r$$


$$s = T\{r\}$$


$$h_s$$

$$s$$



Intensity Transformations

 $f[i, j]$

$$g[i, j] = 2f[i, j] + 1$$

 $g[i, j]$

0	2	2
1	1	1
1	1	2

1	5	5
3	3	3
3	3	5



Intensity Transformations



r	h_r
0	1
1	5
2	3
3	0
4	0
5	0
6	0

$$\begin{aligned}s &= 2r + 1 \\&= 2(1) + 1 \\&= 3\end{aligned}$$

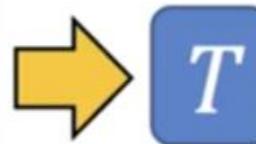
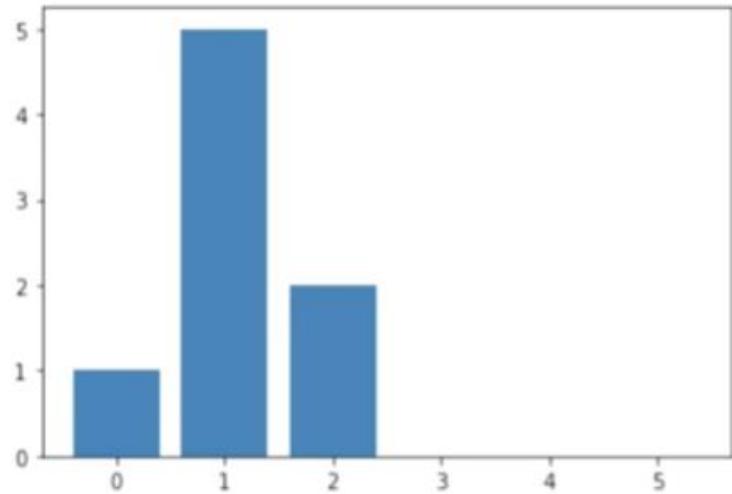
s	h_s
0	0
1	1
2	0
3	5
4	0
5	2
6	0



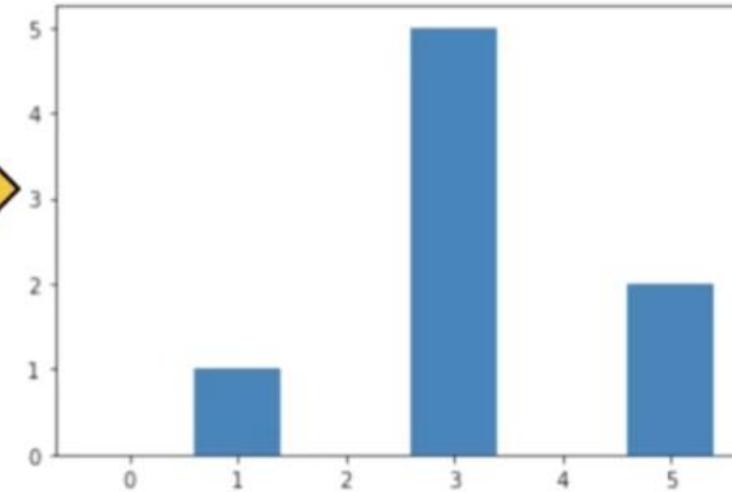
Intensity Transformations



h_r



h_s

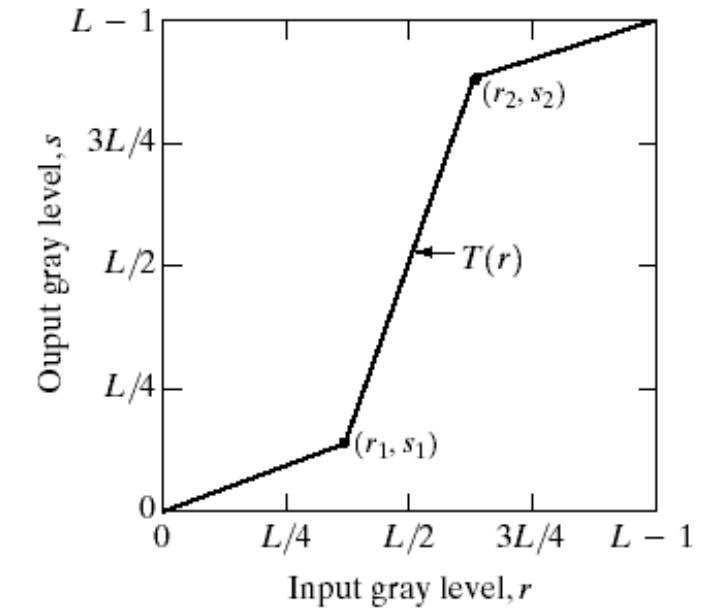
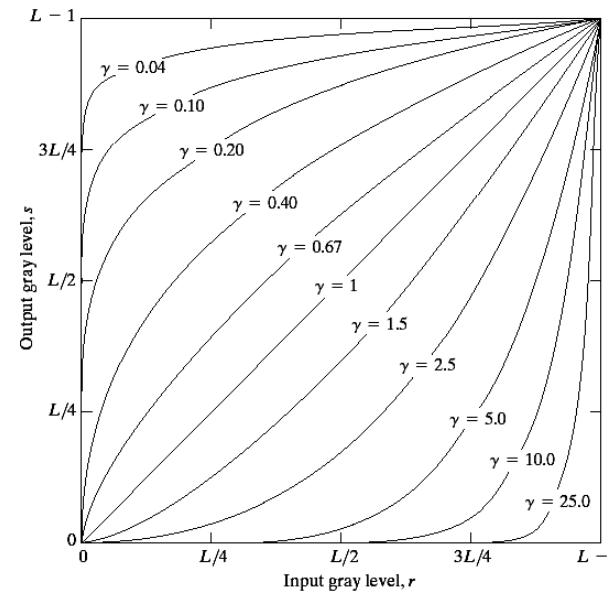
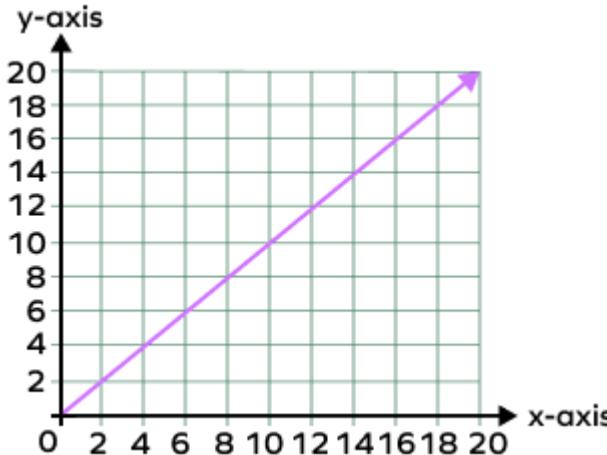




Intensity Transformations



- Linear Transformations (Image Negative)
- Non-Linear Transformations (Log, Inverse Log, Power Law)
- Piecewise-Linear Transformations

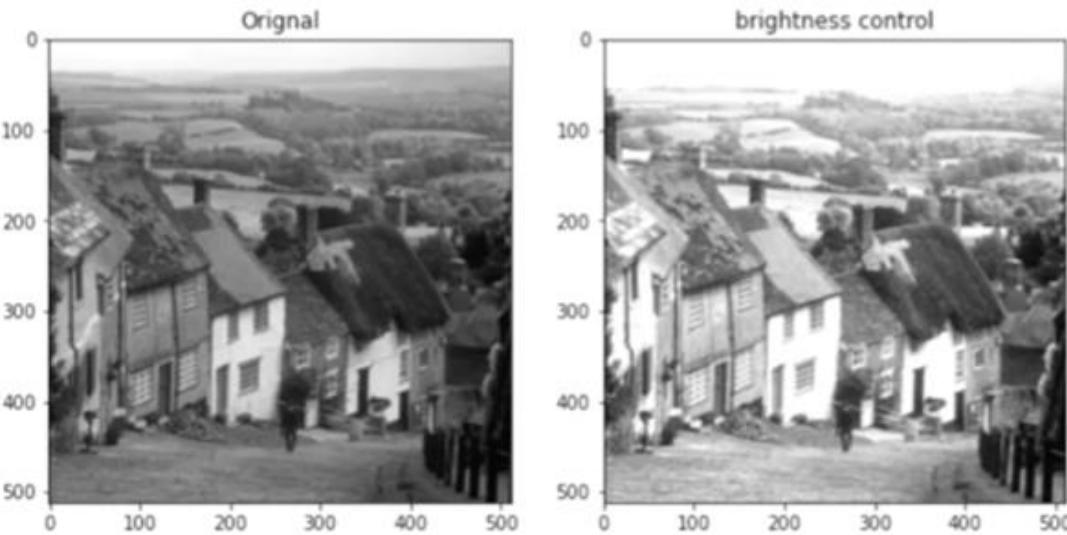




Linear Transformations - Positives

$$g[i,j] = \alpha f[i,j] + \beta$$

```
alpha = 1 # Simple contrast control  
beta = 100 # Simple brightness control  
g[i,j] = 1f[i,j] + 100  
  
new_image = cv2.convertScaleAbs(goldhill, alpha=alpha, beta=beta)
```





Linear Transformations - Negatives

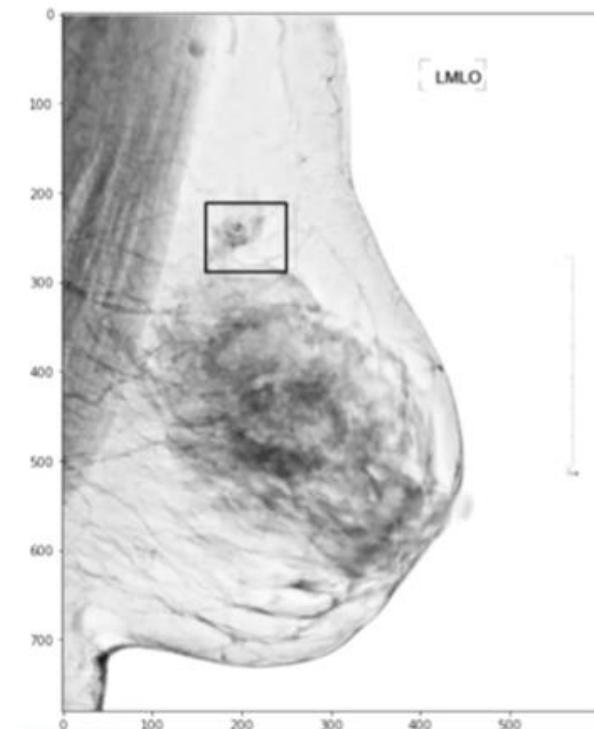


```
image= cv2.imread("mammogram.png",cv2.IMREAD_GRAYSCALE)
```



$$g[i, j] = -1 \cdot f[i, j] + 255$$

```
img_neg=-1* image+255
```



Jian, Wushuai, Xueyan Sun, and Shuqian Luo. "Computer-aided diagnosis of breast microcalcifications based on dual-tree complex wavelet transform." Biomedical engineering online 11.1 (2012): 1-12.



Non-Linear Transformations – Log & Inverse Log



- **Enhances Dark Regions:** The log transform makes dim areas brighter, helping to reveal details that would otherwise be difficult to see.
- **Compresses Bright Pixels:** Very bright pixels have smaller increments after transformation, preventing them from overwhelming the image.
- Opposite is done by Inverse Log.

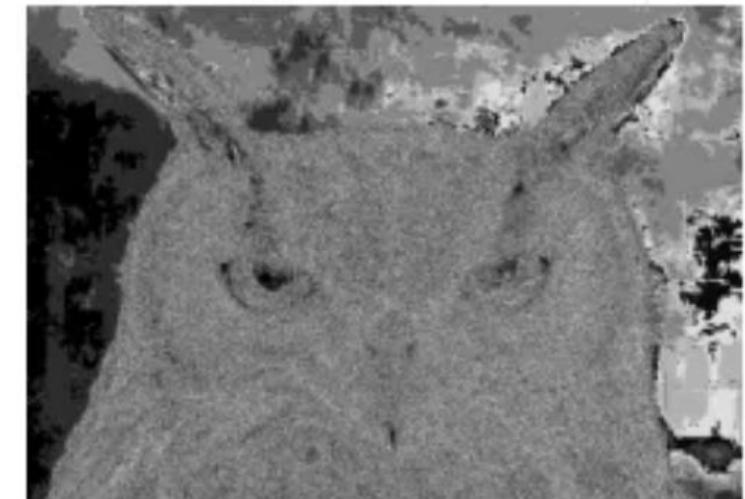
Original Image



log_transformed



inverse_log_transformed





Non-Linear Transformations – Power Law



MR image of
fractured human
spine

Result after
Power law transformation
 $c = 1, \gamma = 0.6$

Result after
Power law transformation
 $c = 1, \gamma = 0.4$

Result after
Power law transformation
 $c = 1, \gamma = 0.3$



Piece-wise Linear Transformation

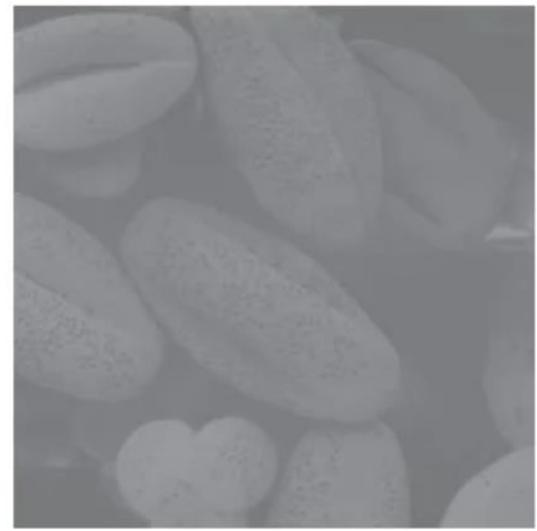
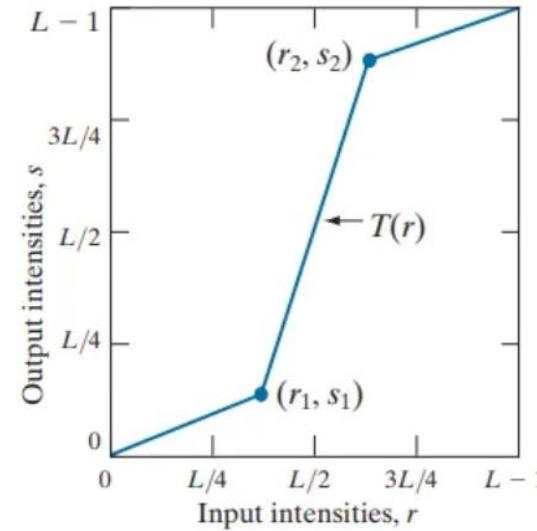


- **Types:**
 - Contrast Stretching
 - Histogram Equalization
 - Thresholding
 - Intensity Level Slicing
 - Bit-plane Slicing

a
b
c
d

FIGURE

Contrast stretching.
(a) Piecewise linear transformation function. (b) A low-contrast electron microscope image of pollen, magnified 700 times.
(c) Result of contrast stretching.
(d) Result of thresholding.
(Original image courtesy of Dr. Roger Heady, Research School of Biological Sciences, Australian National University, Canberra, Australia.)





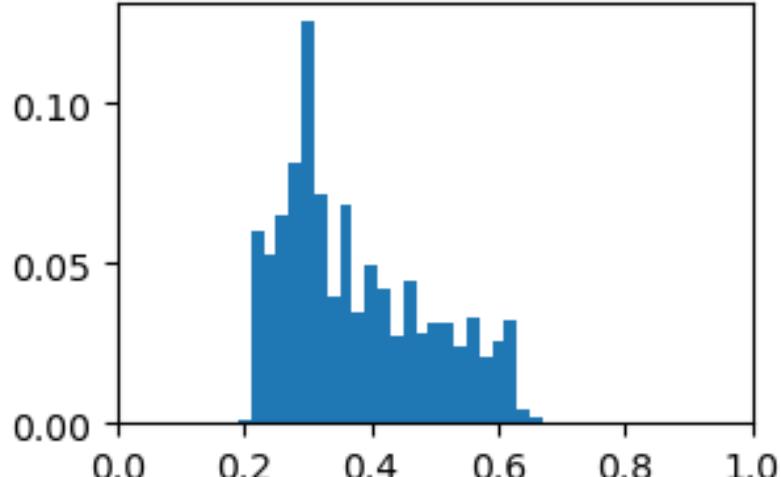
Contrast Stretching



Low contrast orginal



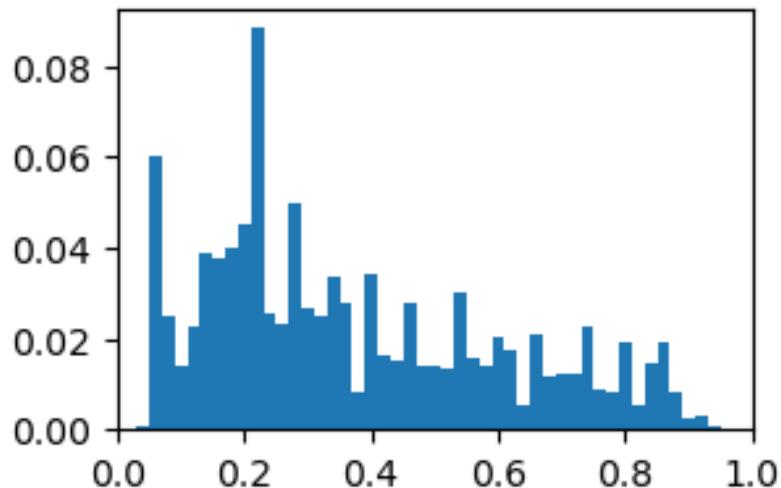
Histogram of low contrast image



Contrast Stretched



Histogram of contrast stretched image



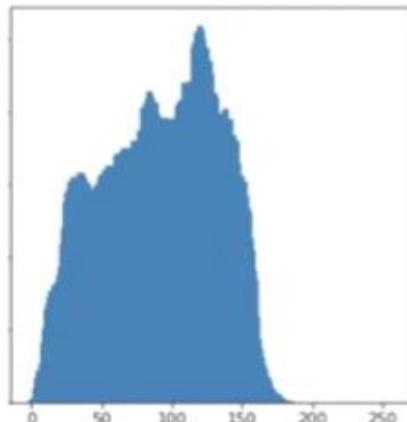
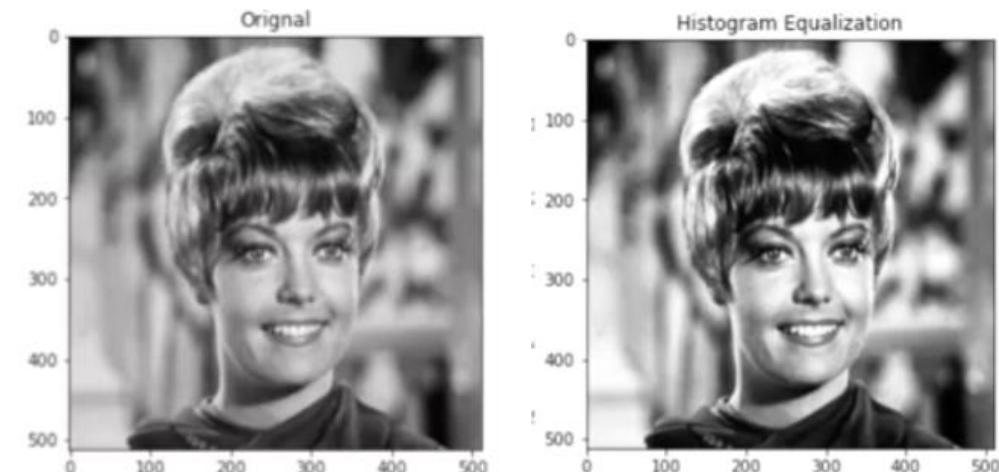


Piecewise-Linear Transformations

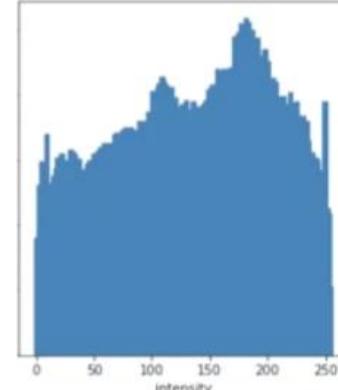
- Histogram Equalization: An algorithm that uses the histogram of images to adjust contrast

```
zelda= cv2.imread("zelda.png",cv2.IMREAD_GRAYSCALE)
```

```
new_image= cv2.equalizeHist(zelda)
```



*Histogram
Equalization*

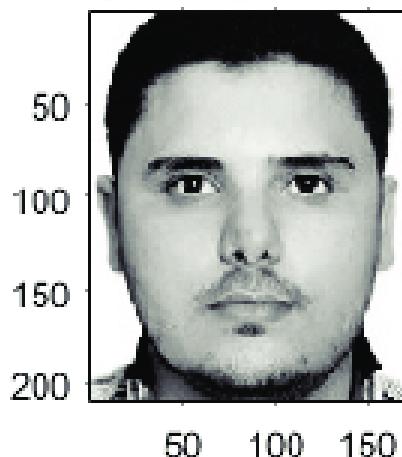




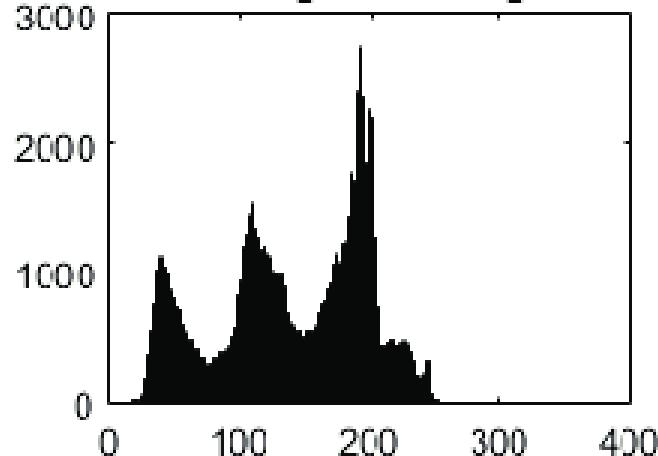
Histogram Equalization



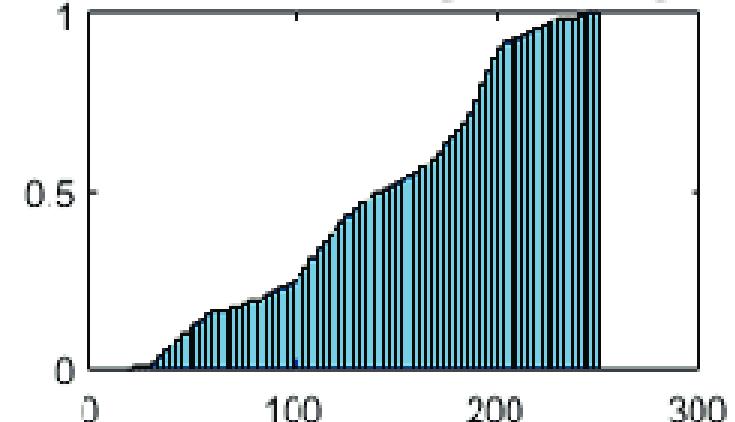
img1: Original Image



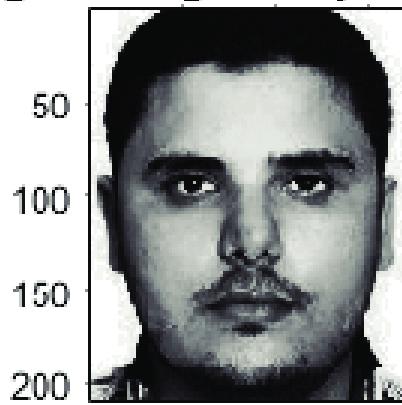
Histogram of img1



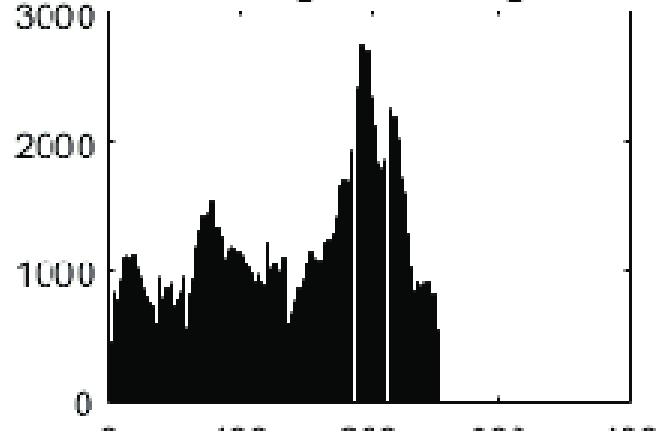
Cumulative Histogram of img1



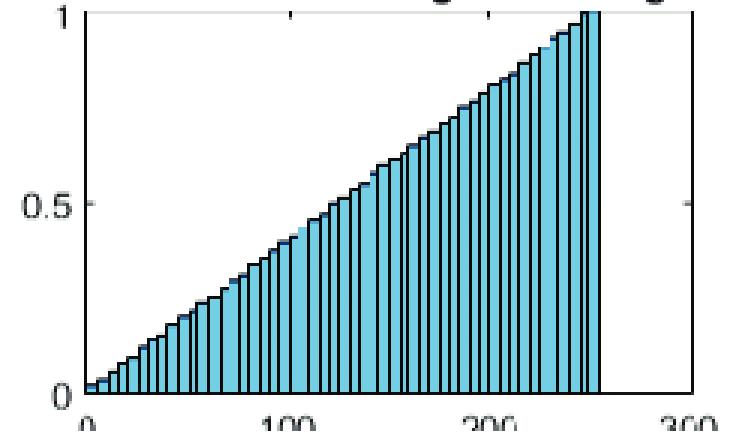
Img2: Histogram Equalization



Histogram of img2



Cumulative Histogram of img2

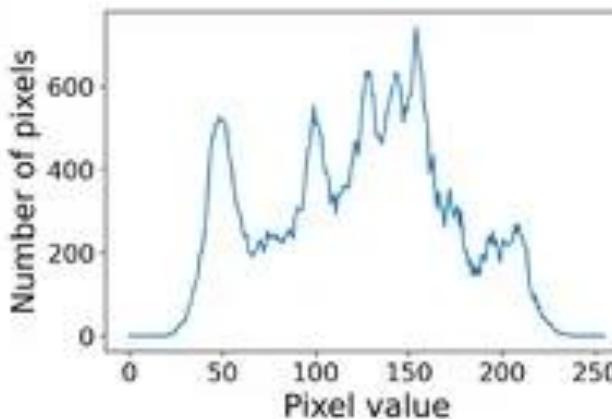




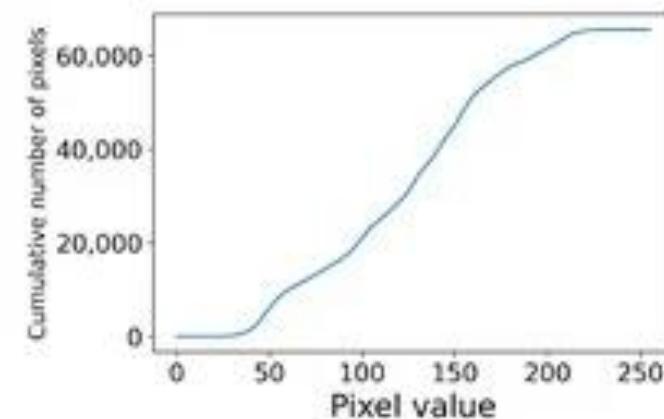
Histogram Equalization



(a) Original image



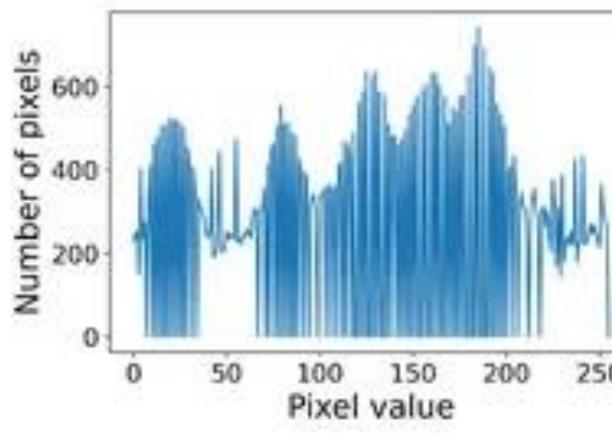
(b) Histogram



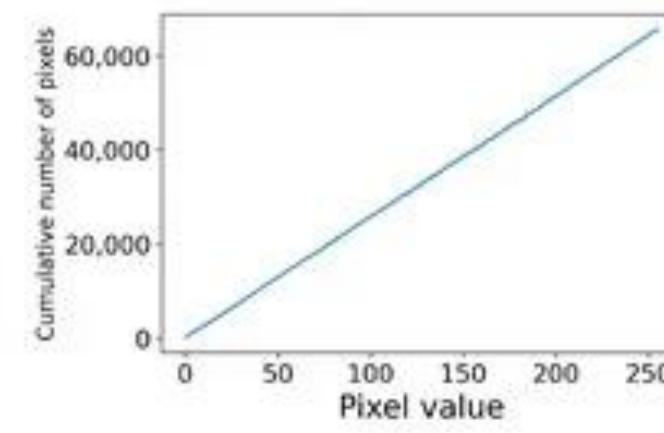
(c) Cumulative histogram



(d) Histogram-equalized image



(e) Histogram



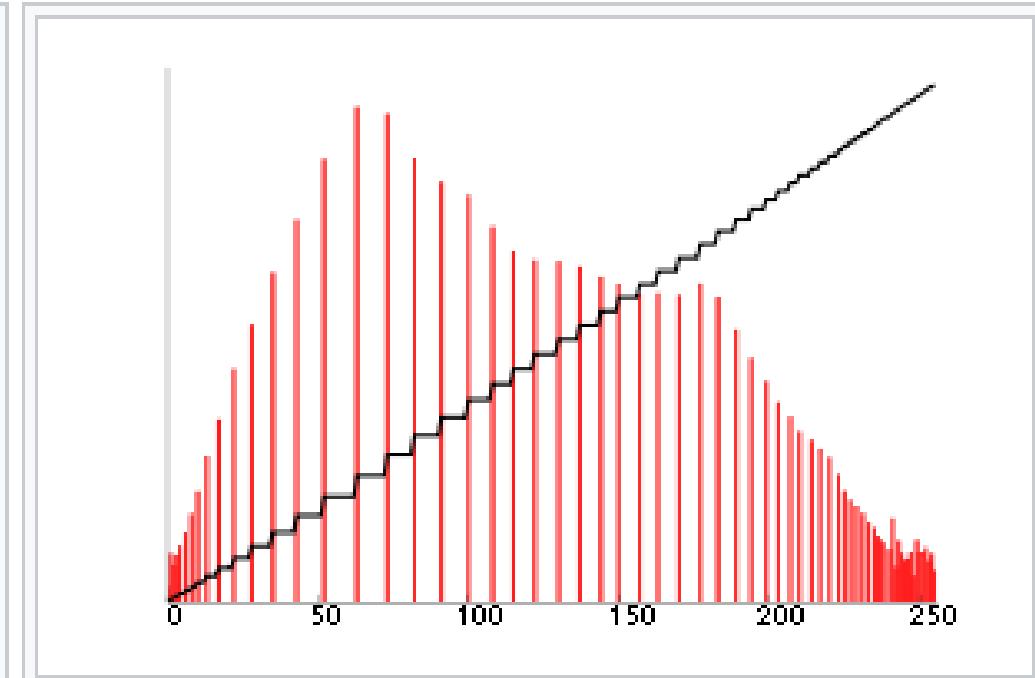
(f) Cumulative histogram



Histogram Equalization



After Histogram Equalization



Corresponding histogram (red) and cumulative histogram (black)





Thresholding and Simple Segmentation



i	input_img		
0	0	2	2
1	1	1	1
2	1	1	2
j	0	1	2

```
for i in range(N):
    for j in range(M):
        if input_img[i,j]> 1:
            image_out[i,j]=255
        else:
            image_out[i,j]=0
```

image_out



Thresholding and Simple Segmentation



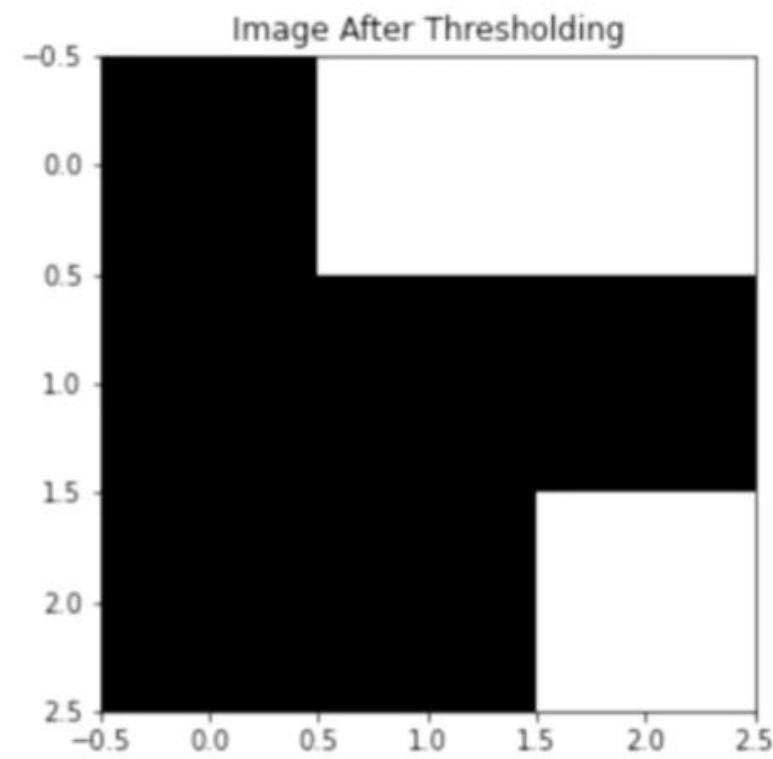
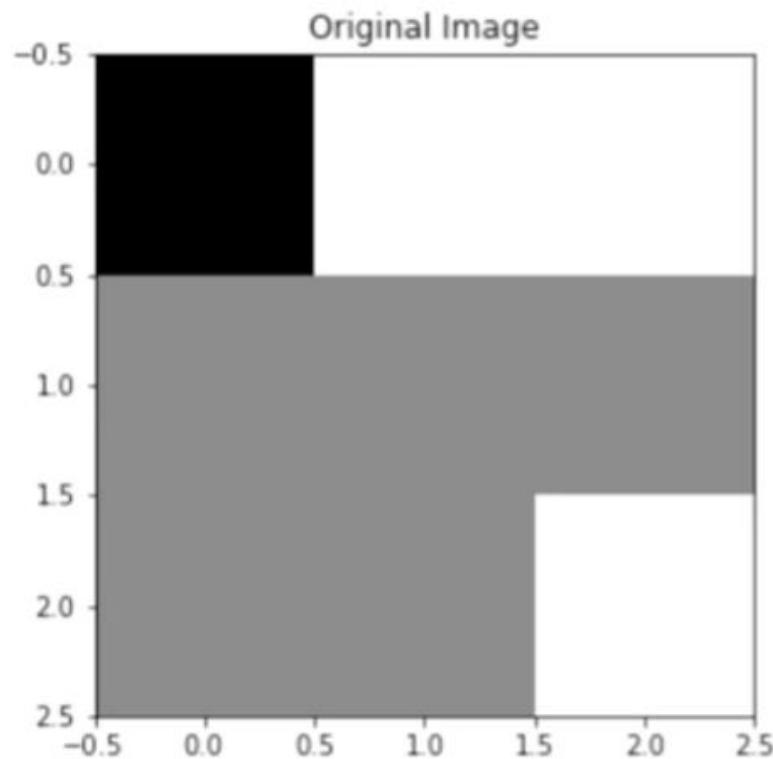
i	input_img		
0	0	2	2
1	1	1	1
2	1	1	2
j	0	1	2

```
for i in range(N):
    for j in range(M):
        if input_img[i,j]> 1:
            image_out[i,j]=2
        else:
            image_out[i,j]=0
```

image_out		
0	255	255
0	0	0



Thresholding and Simple Segmentation





Thresholding and Simple Segmentation

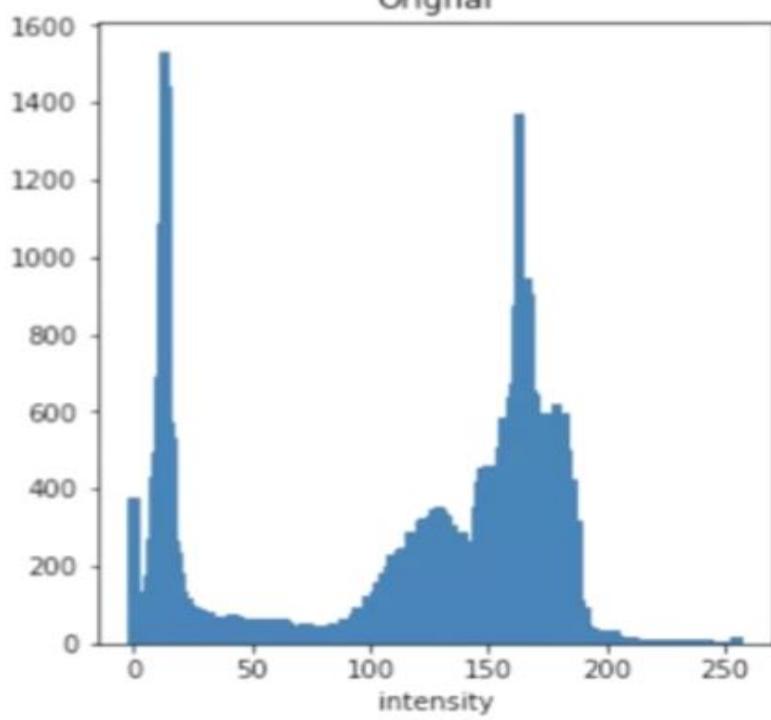


```
image= cv2.imread("cameraman.jpeg",cv2.IMREAD_GRAYSCALE)
```

image



Orignal



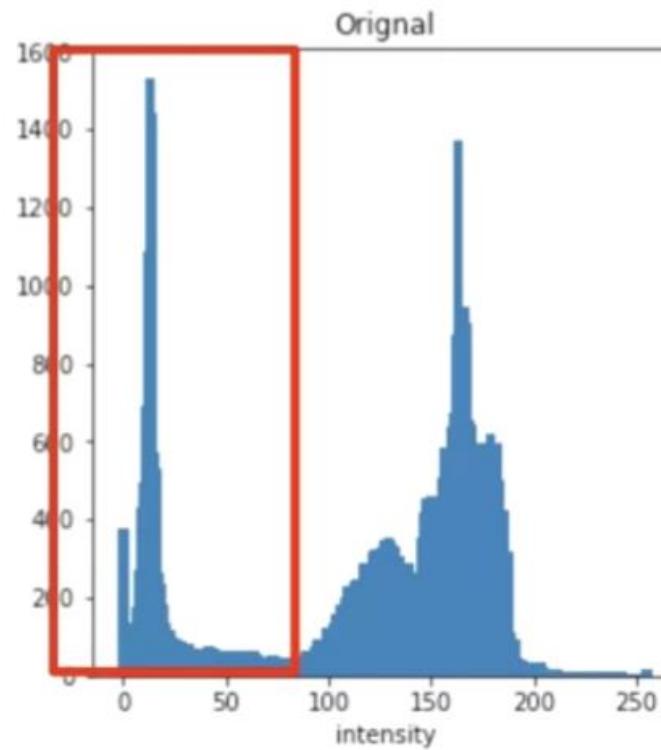
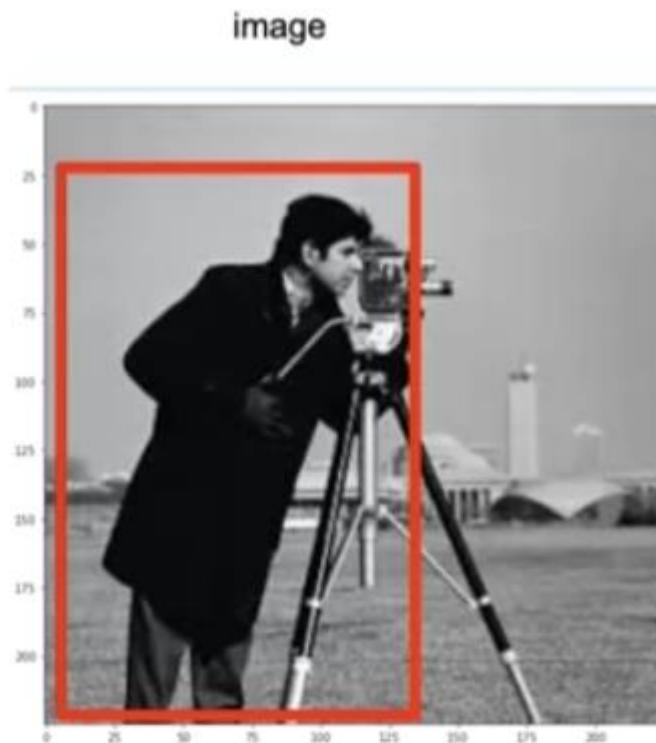


Thresholding and Simple Segmentation

```
image= cv2.imread("cameraman.jpeg",cv2.IMREAD_GRAYSCALE)
```

```
max_value=255
```

```
threshold=87
```





Thresholding and Simple Segmentation

```
image= cv2.imread("cameraman.jpeg",cv2.IMREAD_GRAYSCALE)
```

```
max_value=255
```

```
threshold=87
```

```
ret, new_image = cv2.threshold(image,threshold,max_value,cv2.THRESH_BINARY)
```

image



Original

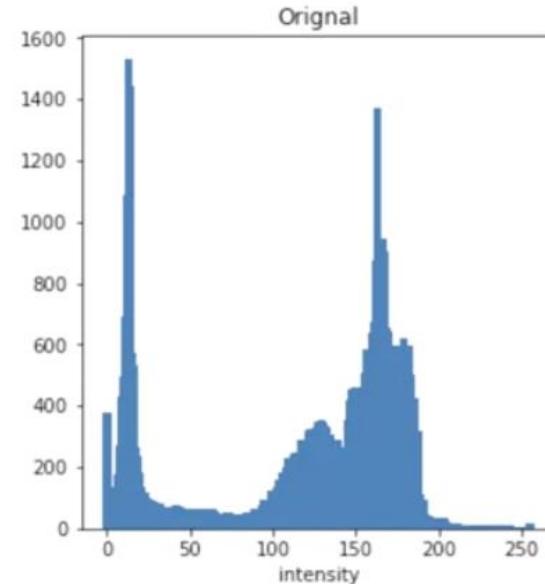


Image After Thresholding

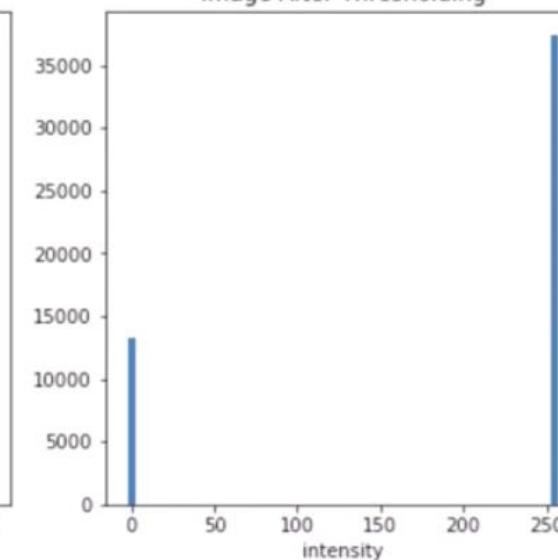
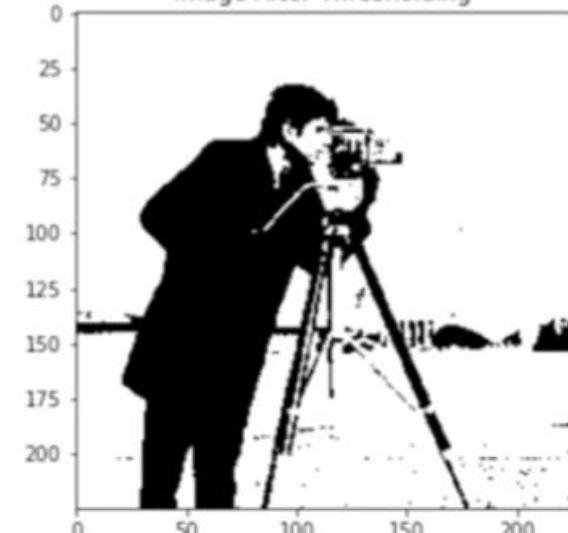


Image After Thresholding



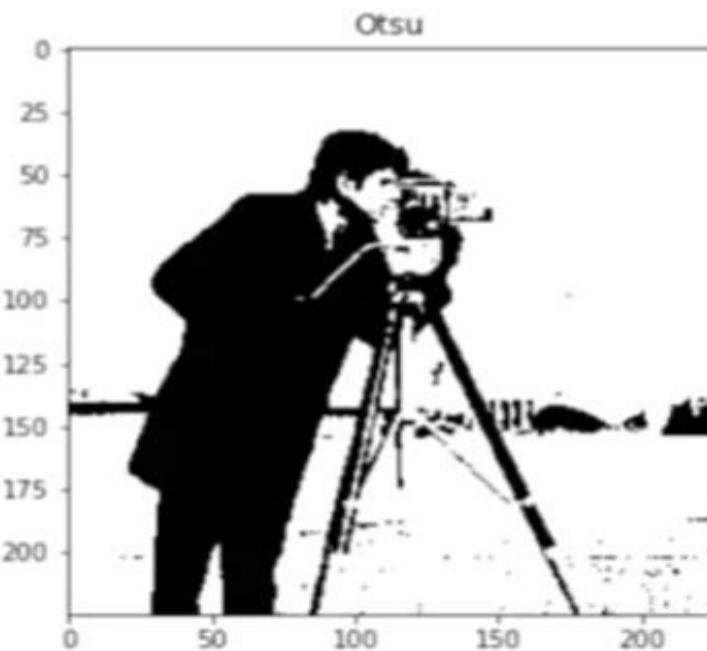


Thresholding and Simple Segmentation



```
ret, otsu = cv2.threshold(image,0,255,cv2.THRESH_OTSU)
```

ret:88.0





Intensity Level Slicing

Highlighting a specific range of gray levels in an image

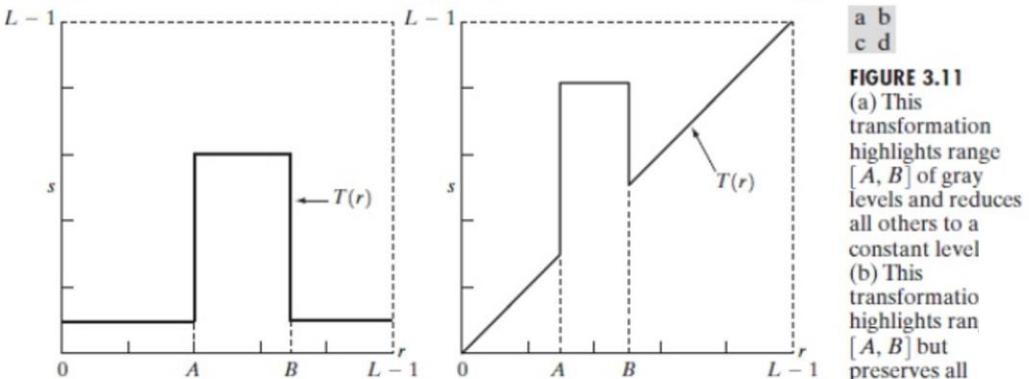


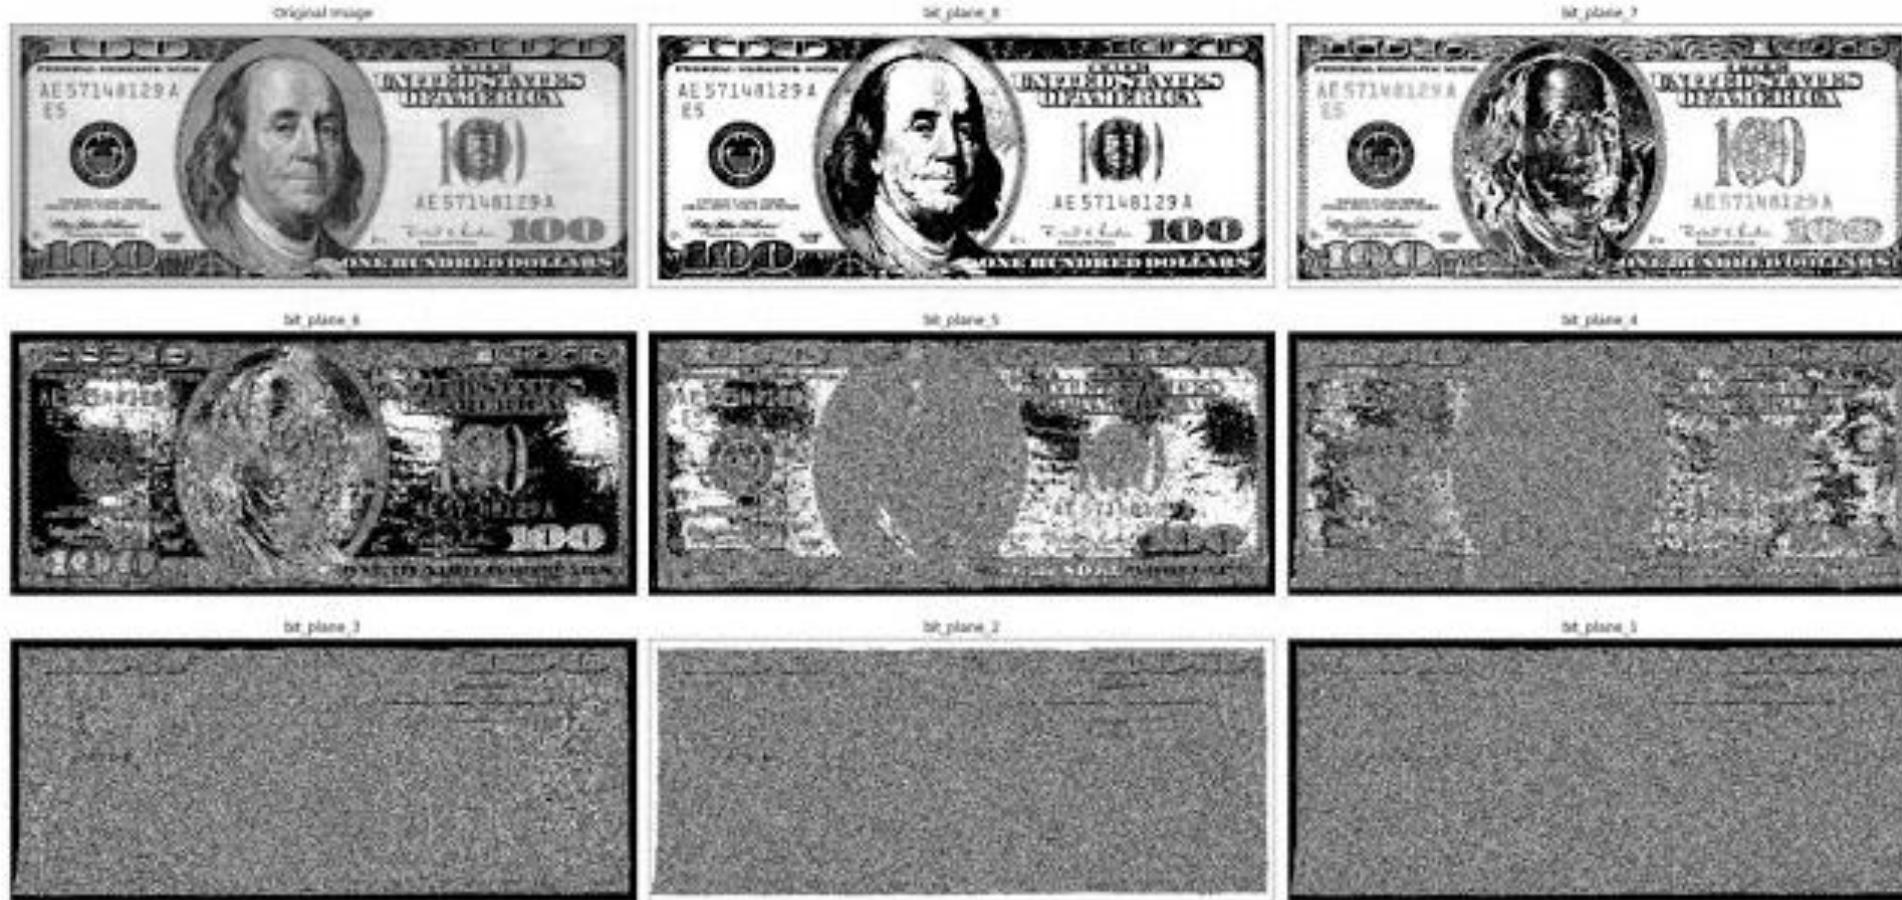
FIGURE 3.11
(a) This transformation highlights range $[A, B]$ of gray levels and reduces all others to a constant level
(b) This transformation highlights range $[A, B]$ but preserves all other levels



FIGURE 3.12 (a) Aortic angiogram. (b) Result of using a slicing transformation of the type illustrated in Fig. 3.11(a), with the range of intensities of interest selected in the upper end of the gray scale. (c) Result of using the transformation in Fig. 3.11(b), with the selected area set to black, so that grays in the area of the blood vessels and kidneys were preserved. (Original image courtesy of Dr. Thomas R. Gest, University of Michigan Medical School.)



Bit-plane Slicing

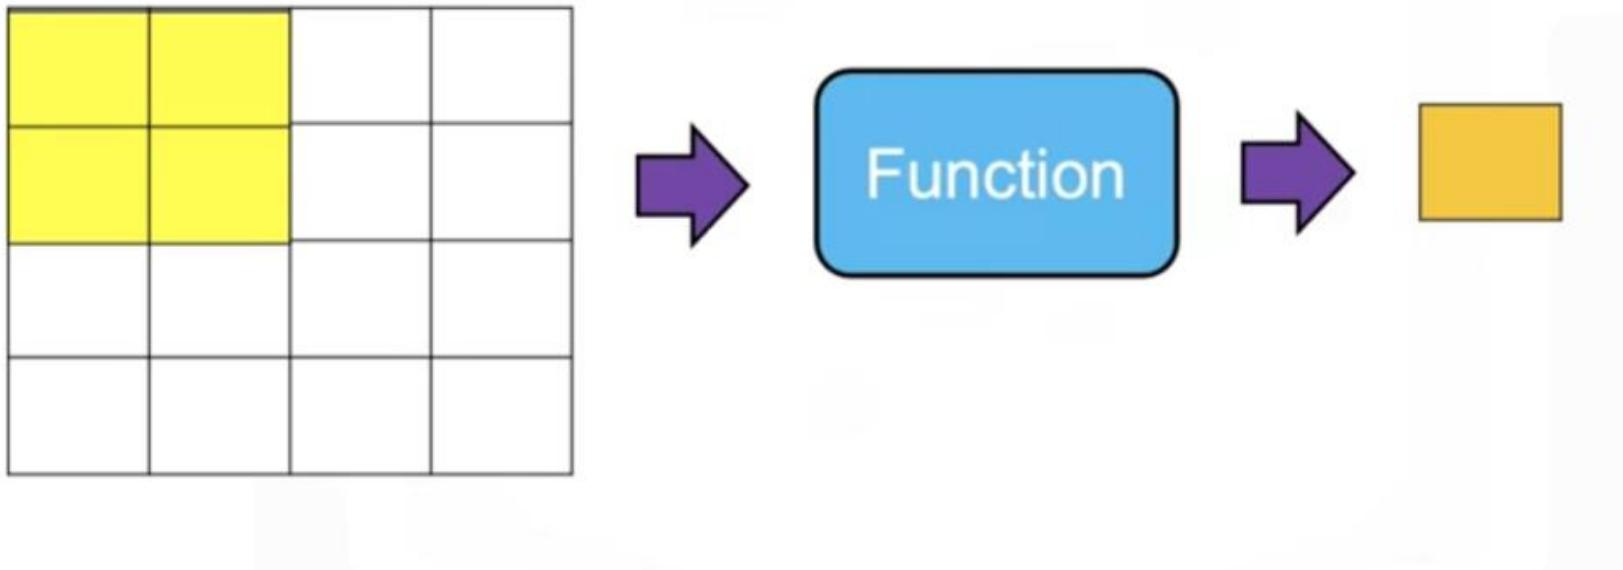




Spatial Operations



- Spatial operations consist of a neighborhood, we apply a function on the neighborhood and obtain the result.

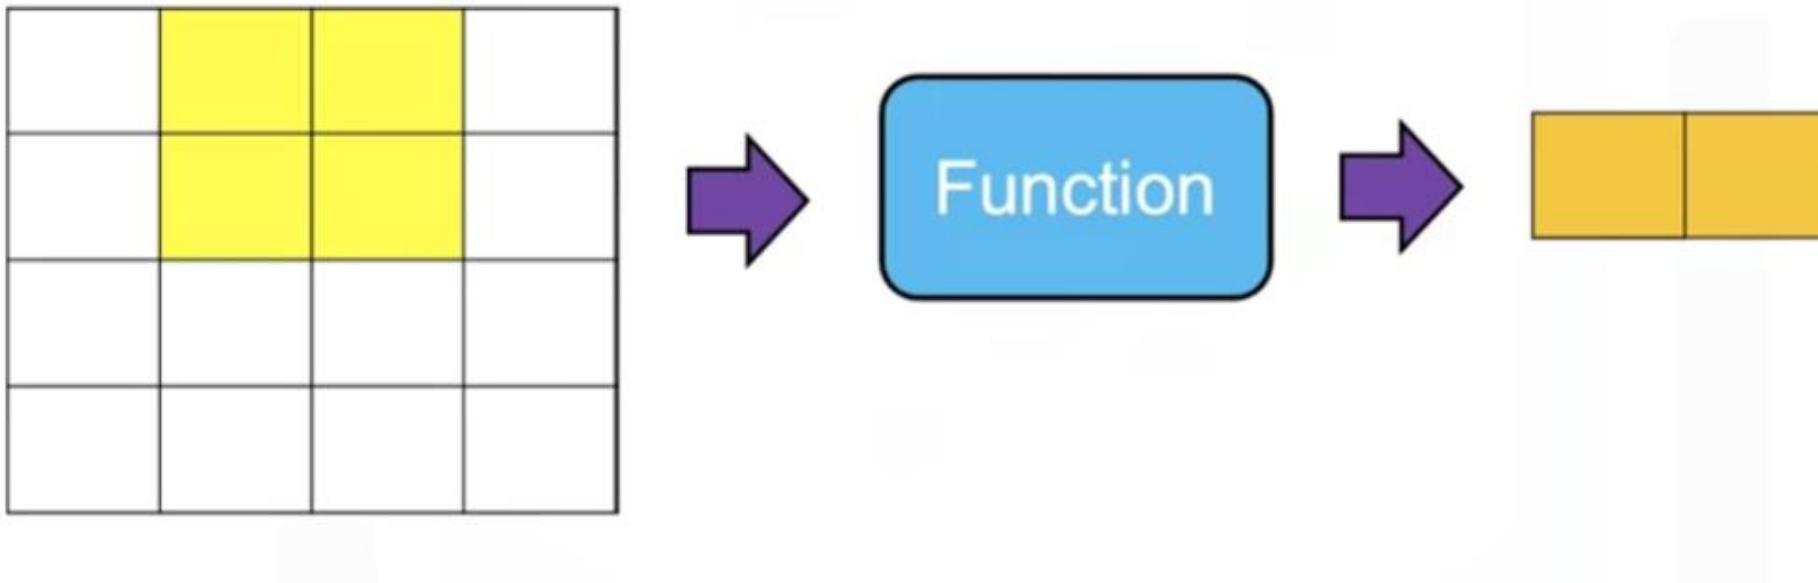




Spatial Operations



- We then shift the neighborhood using the sliding window approach and obtain the result for the next pixel in the resultant image.

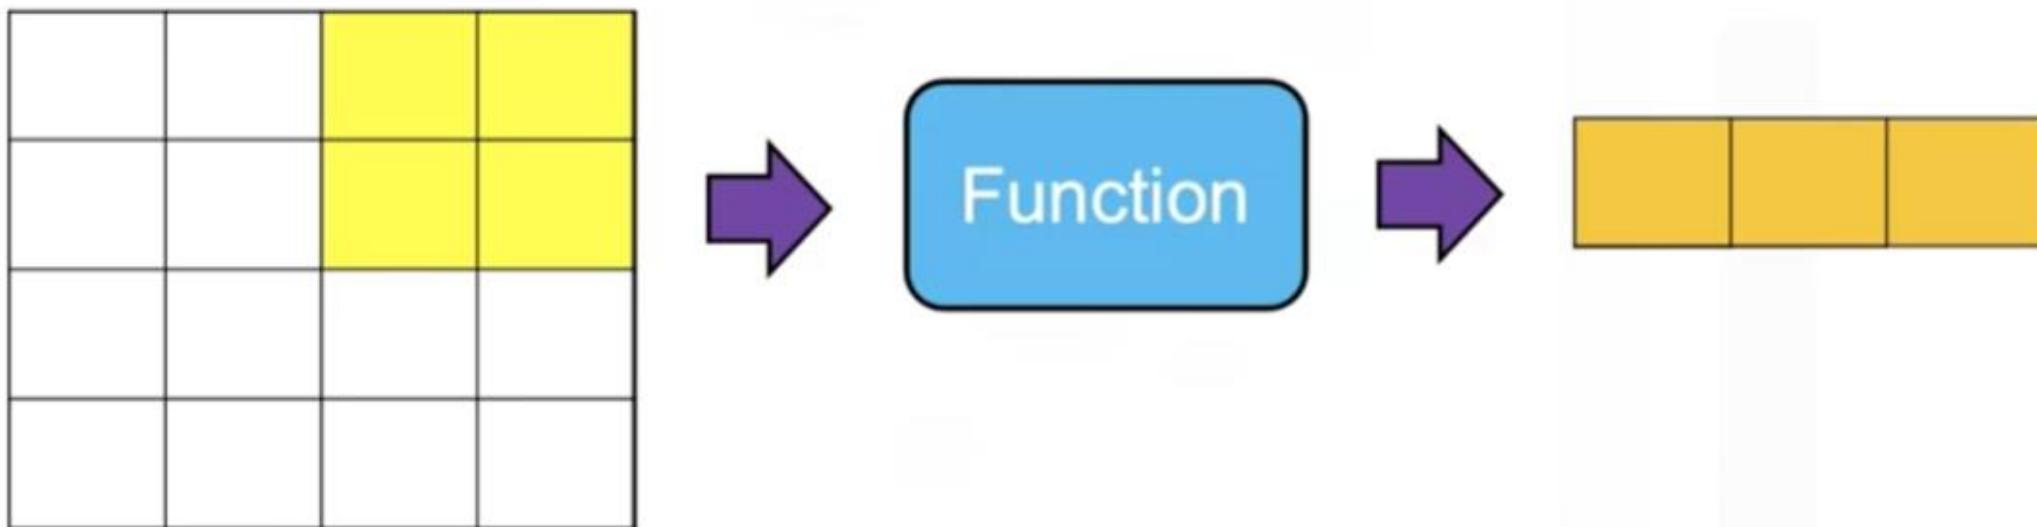




Spatial Operations



- We then shift the neighborhood using the sliding window approach and obtain the result for the next pixel in the resultant image.

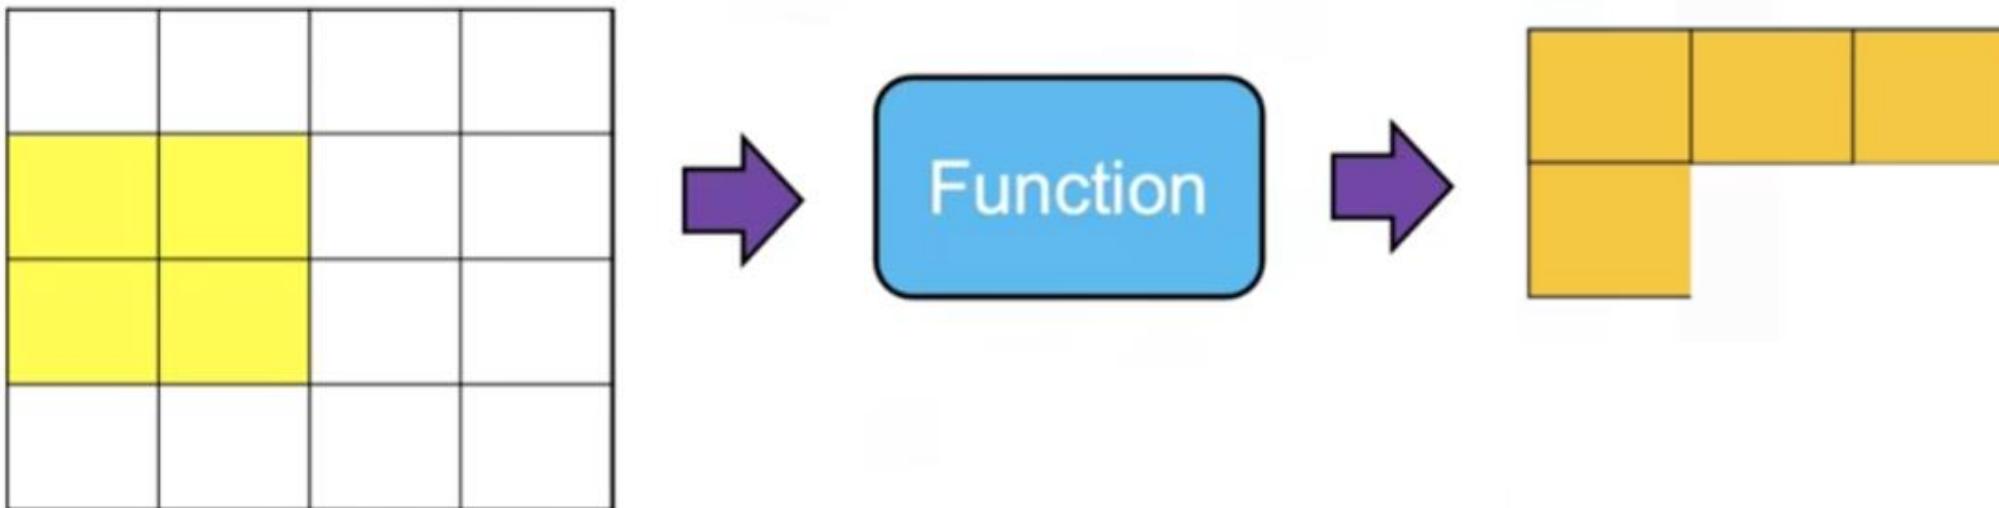




Spatial Operations



- We then shift the neighborhood using the sliding window approach and obtain the result for the next pixel in the resultant image.

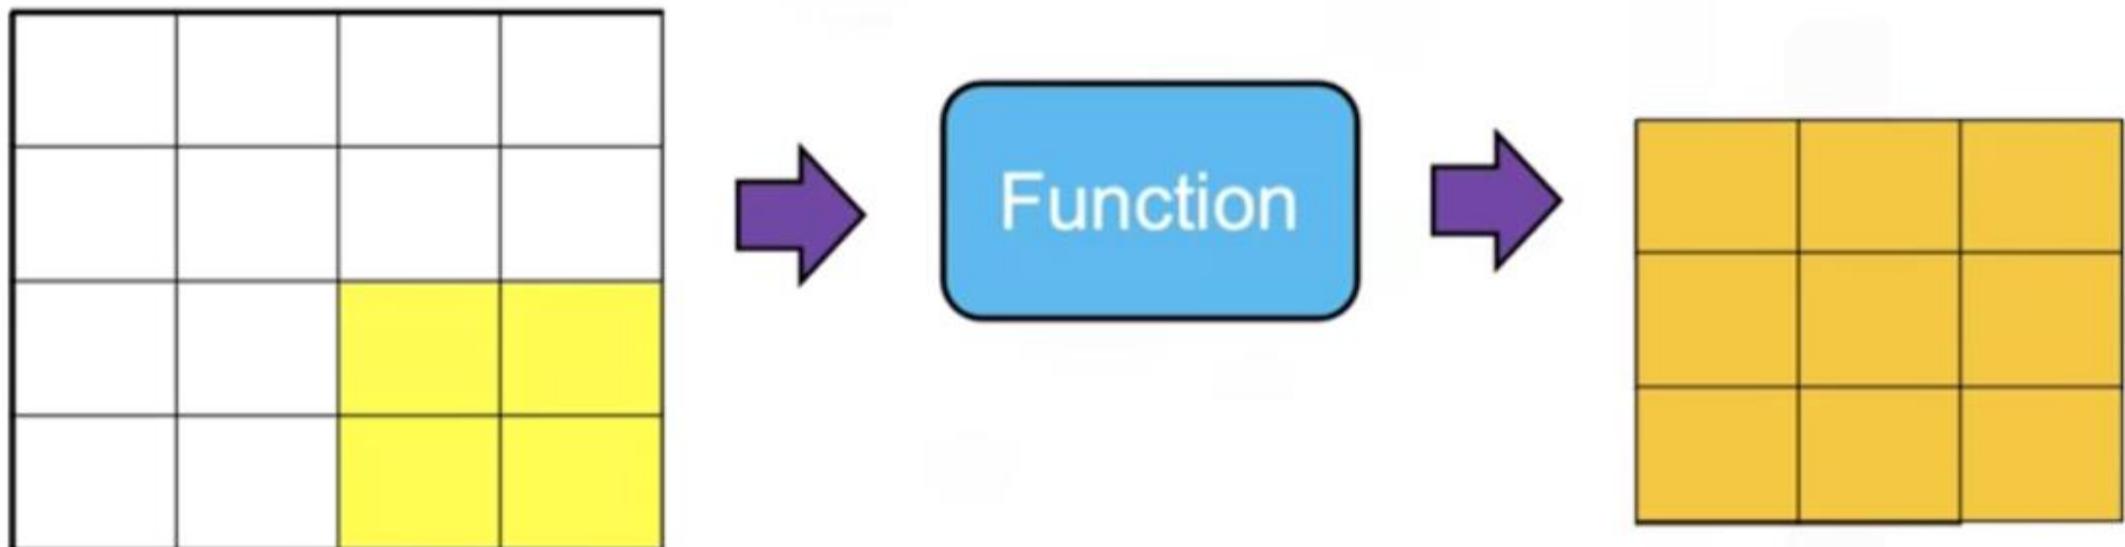




Spatial Operations



- The result is a new image that has enhanced characteristics. These operations take advantage of the spatial patterns in the image.

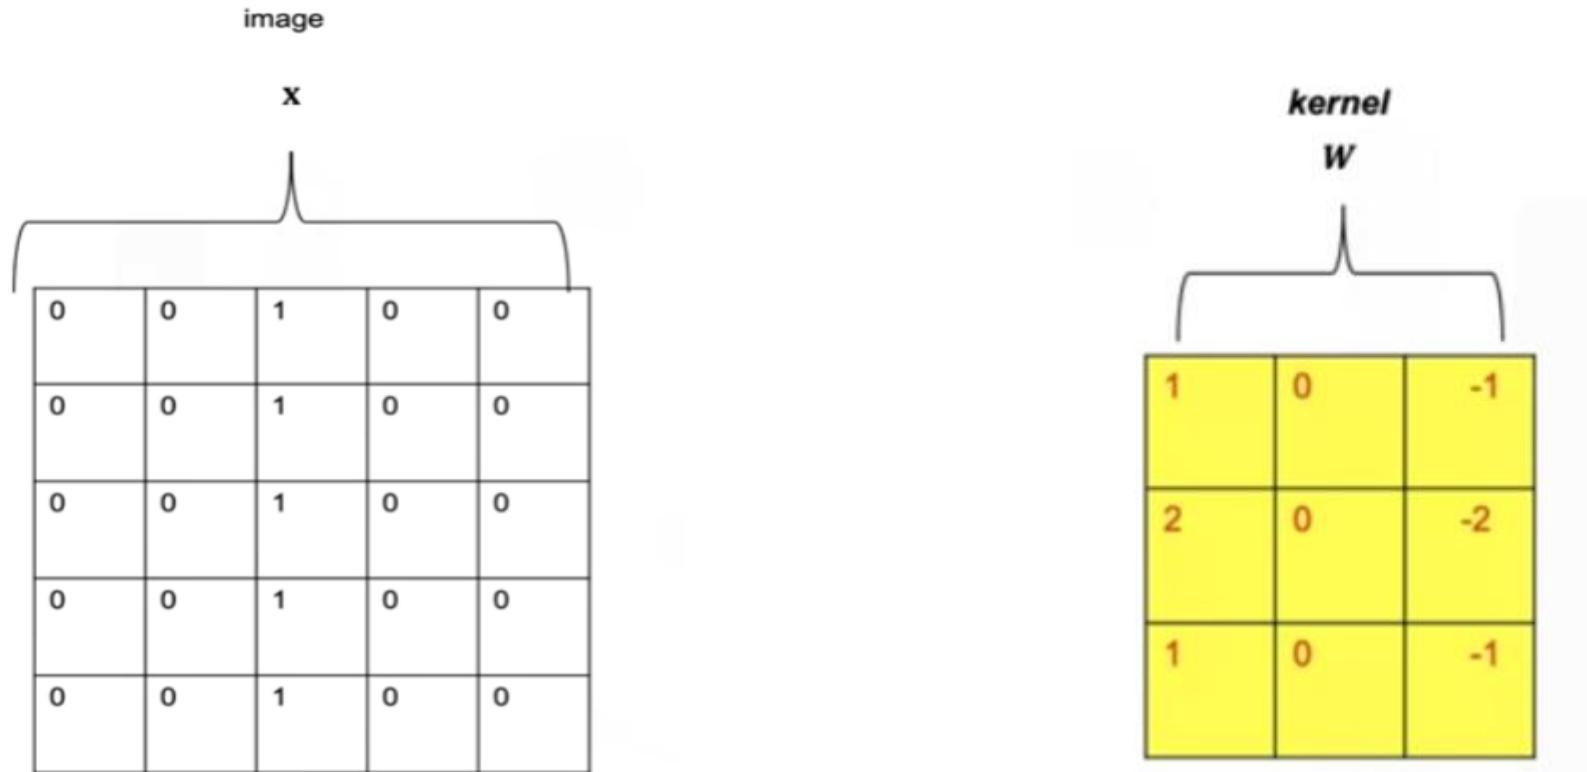




Kernels



A kernel is a filter that performs a specific task on the neighboring pixels determined by the selected window size.





Convolution – Linear Filtering

Convolution is a linear operation that multiplies the corresponding values of the filter with the pixels in the image and sums the prod

X

$$\mathbf{Z} = \mathbf{W} * \mathbf{X}$$

0	1	0	0	1	-1	0	0
0	2	0	0	1	-2	0	0
0	1	0	0	1	-1	0	0
0	0	0	1	0	0	0	0
0	0	1	0	0	0	0	0

1	0	-1
2	0	-2
1	0	-1



Convolution



Convolution is a linear filter that multiplies the corresponding values of the filter with the pixels in the image and sums the products.

X

0	1	0	0	1	-1	0	0
0	2	0	0	1	-2	0	0
0	1	0	0	1	-1	0	0
0	0	1		0		0	
0	0	1		0		0	

$$0 \times 1 + 0 \times 0 + 1 \times -1$$

$$\mathbf{Z} = \mathbf{W} * \mathbf{X}$$



Convolution



Convolution is a linear filter that multiplies the corresponding values of the filter with the pixels in the image and sums the products.

X

0	1	0	0	1	-1	0	0
0	2	0	0	1	-2	0	0
0	1	0	0	1	-1	0	0
0	0	1		0	0		
0	0	1		0	0		

$$0x1+0x0+1x-1$$

+

$$0x2+0x0+1x-2$$

$$\mathbf{Z} = \mathbf{W} * \mathbf{X}$$



Convolution



Convolution is a linear filter that multiplies the corresponding values of the filter with the pixels in the image and sums the products.

$$\mathbf{Z} = \mathbf{W} * \mathbf{X}$$

X

0	1	0	0	1	-1	0	0
0	2	0	0	1	-2	0	0
0	1	0	0	1	-1	0	0
0	0	1		0	0		
0	0	1		0	0		

$$0x1+0x0+1x-1$$

+

$$0x2+0x0+1x-2$$

+

$$0x1+0x0+1x-1$$



Convolution



Convolution is a linear filter that multiplies the corresponding values of the filter with the pixels in the image and sums the products.

$$Z = W * X$$

X

0	1	0	0	1	-1	0	0
0	2	0	0	1	-2	0	0
0	1	0	0	1	-1	0	0
0	0	1		0	0	0	0
0	0	1		0	0	0	0

Z

$$\begin{array}{l} 0x1+0x0+1x-1 \\ + \\ 0x2+0x0+1x-2 \\ + \\ 0x1+0x0+1x-1 \end{array}$$

-1
+
-2
+
-1
4

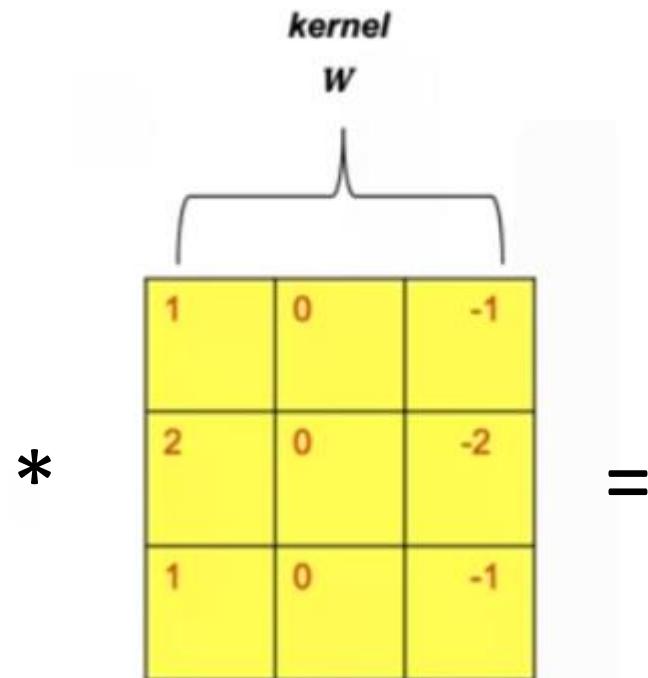
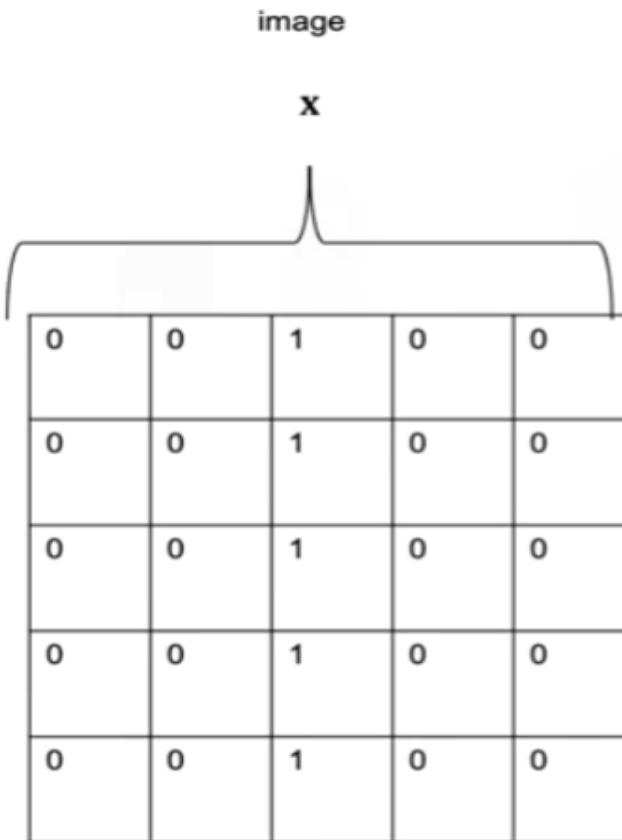
-4



Convolution



$$\mathbf{Z} = \mathbf{W} * \mathbf{X}$$



Resultant Image is
Smaller in size

\mathbf{Z}

-4	0	4
-4	0	4
-4	0	4



Different Types of Kernels/Filters

- Using the kernel operations different types of kernels can create different effect when applied on the image.
- In general, some of them provide smoothing/noise reduction and some of them are for edge detection/enhancement.
- They apply on the original image using convolution operation.

Kernel Type	Matrix Example	Application
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	Original image reproduction
Sharpening	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	Detail and edge enhancement



Different Types of Kernels/Filter

Kernel Type	Matrix Example	Application
Box Blur	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	Basic blur and noise reduction
Gaussian Blur	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	Smoother blur effect for noise reduction
Sobel (Edge)	$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$	Edge detection (horizontal/vertical)
Laplacian (Edge)	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	Second derivative edge detection
Emboss	$\begin{bmatrix} -2 & -1 & 0 \\ -1 & 1 & 1 \\ 0 & 1 & 2 \end{bmatrix}$	3D or embossed appearance
Edge Enhancement	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	Intense edge detail highlighting



Different Types of Kernels/Filters



- Smoothing/Noise Reduction Filters are also called as Low-pass filter. They can be linear or non-linear.
- Sharpening/Edge Detection filters are also called as High-pass filter. They are usually linear.





Different Types of Kernels/Filters



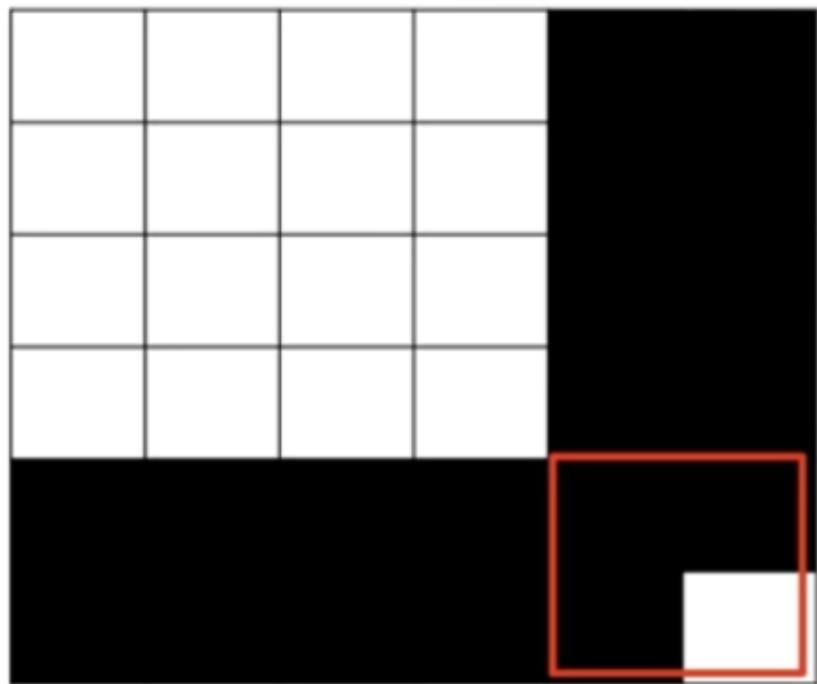
- **Low-Pass Linear Filter:** Blurring an image to reduce noise using a Gaussian filter.
 - **High-Pass Linear Filter:** Detecting edges using a Laplacian filter.
 - **Low-Pass Non-Linear Filter:** Removing salt-and-pepper noise with a median filter while keeping edges intact.
-
- For understanding the effect of different types of kernels/filters, follow this link: <https://setosa.io/ev/image-kernels/>
 - For implementation of different types of kernels/filters in simple Pythons, follow this link: <https://medium.com/swlh/image-processing-with-python-convolutional-filters-and-kernels-b9884d91a8fd>



Low Pass Filters



Low pass filters are used to smooth (blur) the image to remove noise.



$$\text{kernel} \\ * \begin{array}{|c|c|} \hline 1/4 & 1/4 \\ \hline 1/4 & 1/4 \\ \hline \end{array} =$$

Mean (Averaging) Filter



Low Pass Filters



Low pass filters are used to smooth (blur) the image to remove noise.

255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
0	0	0	0	0	0
0	0	0	0	0	255

$$\begin{matrix} & \text{kernel} \\ * & \begin{matrix} 1/4 & 1/4 \\ 1/4 & 1/4 \end{matrix} \\ = & \end{matrix}$$

Mean (Averaging) Filter



Low Pass Filters



Low pass filters are used to smooth (blur) the image to remove noise.

255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
0	0	0	0	0	0
0	0	0	0	0	255

$$\begin{matrix} & \text{kernel} \\ * & \begin{array}{|c|c|} \hline & 1/4 & 1/4 \\ \hline 1/4 & & 1/4 \\ \hline \end{array} \\ = & \end{matrix}$$

255	255	255			
255	255	255			
255	255	255			

Mean (Averaging) Filter



Low Pass Filters



Low pass filters are used to smooth (blur) the image to remove noise.

255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
0	0	0	0	0	0
0	0	0	0	0	255

$$\begin{matrix} & \text{kernel} \\ & \begin{array}{|c|c|} \hline 1/4 & 1/4 \\ \hline 1/4 & 1/4 \\ \hline \end{array} \\ * & \end{matrix} =$$

255	255	255	255	128	0
255	255	255	255	128	0
255	255	255	255	128	0
255	255	255	255	128	0
128	128	128	128	64	0
0	0	0	0	0	64

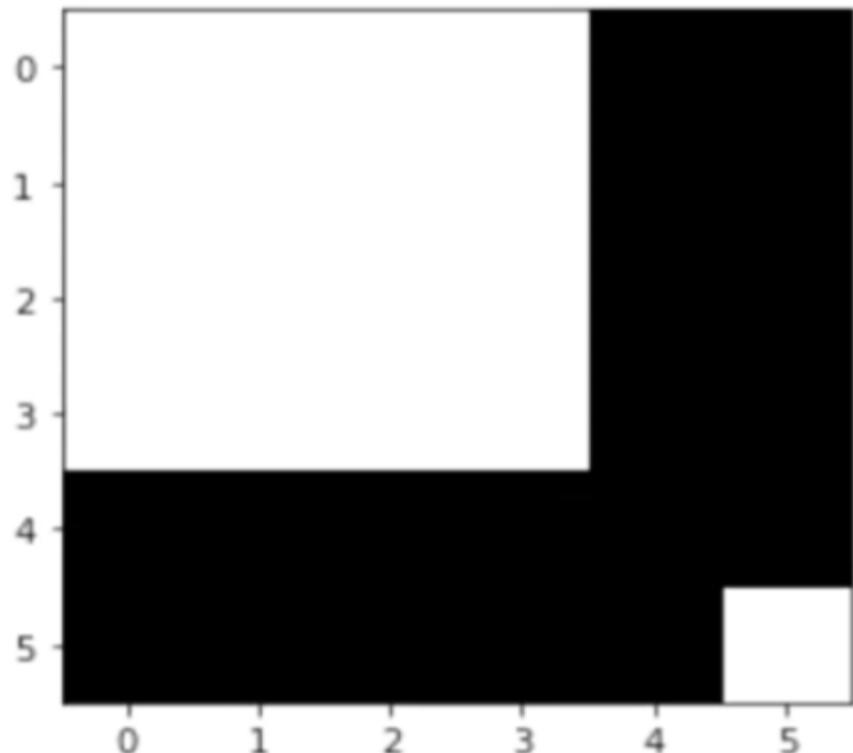
Mean (Averaging) Filter



Low Pass Filters

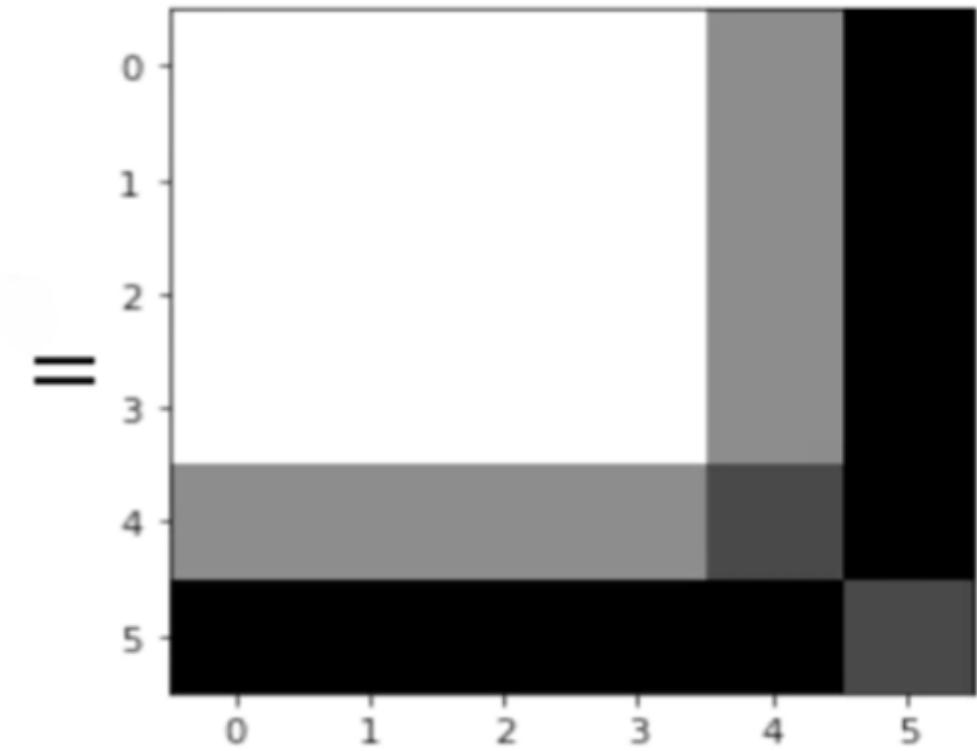


Low pass filters are used to smooth (blur) the image to remove noise.



kernel

$$\begin{matrix} * & \begin{matrix} 1/4 & 1/4 \\ 1/4 & 1/4 \end{matrix} \\ \hline \end{matrix}$$



Mean (Averaging) Filter

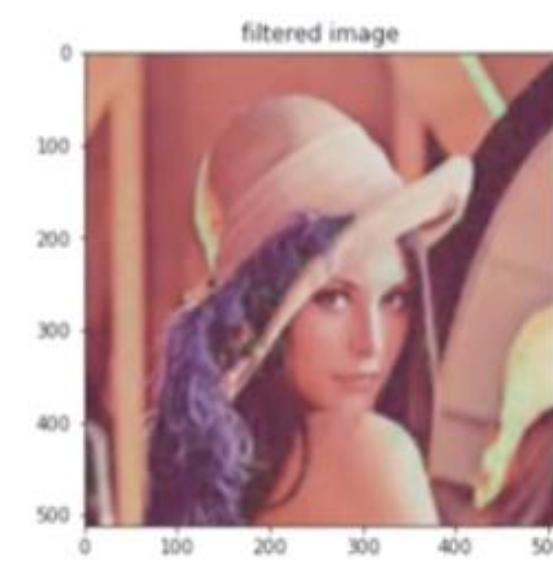
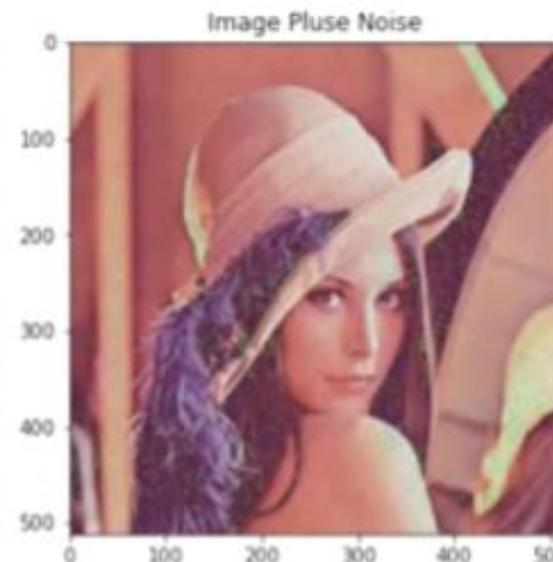


Implementation – Mean Filtering

```
new_image= image+Noise
```

```
kernel = np.ones((6,6))/36
```

```
image_filtered=cv2.filter2D(src=new_image ,ddepth=-1,kernel=kernel)
```

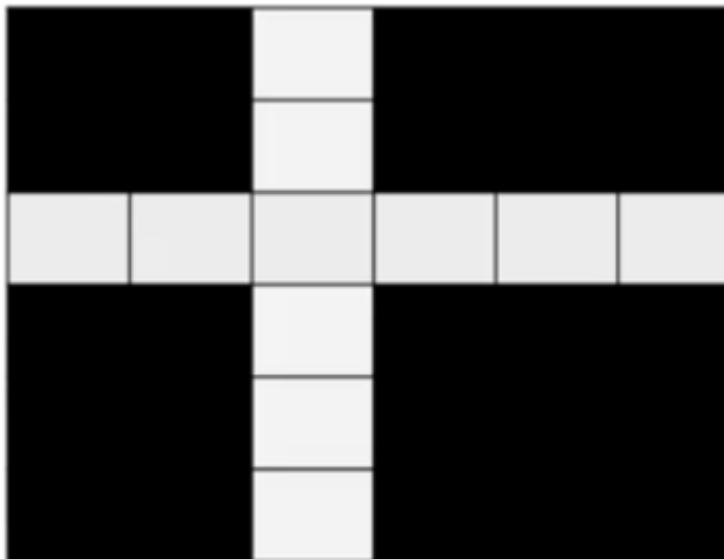




Edge Detection Filters



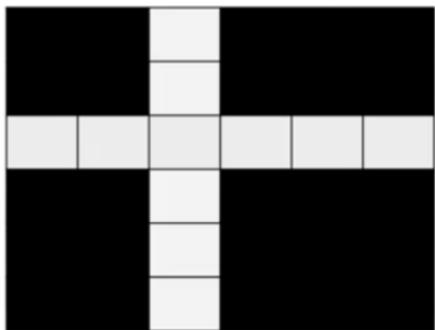
Edges are the points in an image where the intensity levels (image brightness) changes sharply. Edge detection uses methods to approximate derivatives and gradients to identify these points.



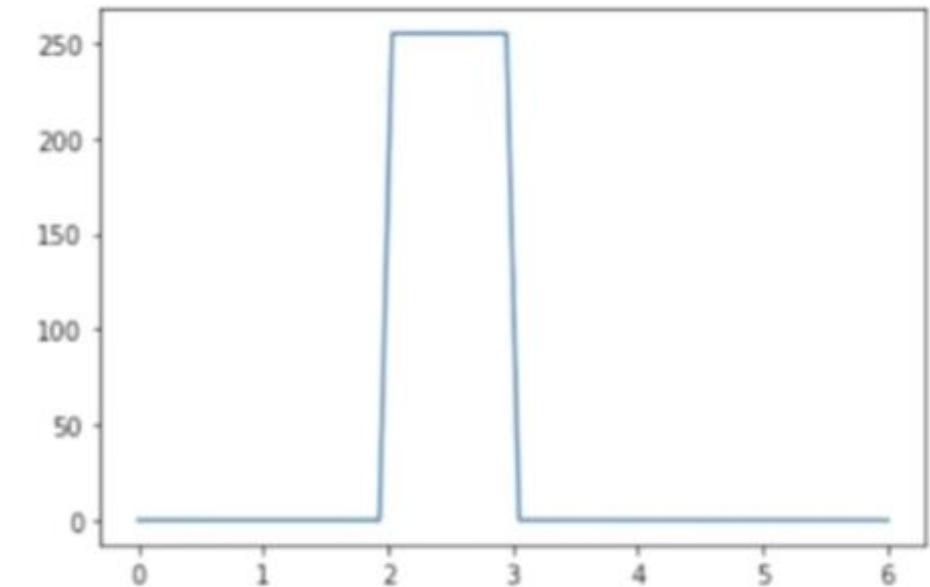


Edge Detection Filters

Edges are the points in an image where the intensity levels (image brightness) changes sharply. Edge detection uses methods to approximate derivatives and gradients to identify these points.



0	0	255	0	0	0
0	0	255	0	0	0
255	255	255	255	255	255
0	0	255	0	0	0
0	0	255	0	0	0
0	0	255	0	0	0





Edge Detection Filters



0	1	2	3	4	5
0	-0	255	0	0	0
0	-0	255	0	0	0
255	255	255	255	255	255
0	-0	255	0	0	0
0	-0	255	0	0	0
0	-0	255	0	0	0

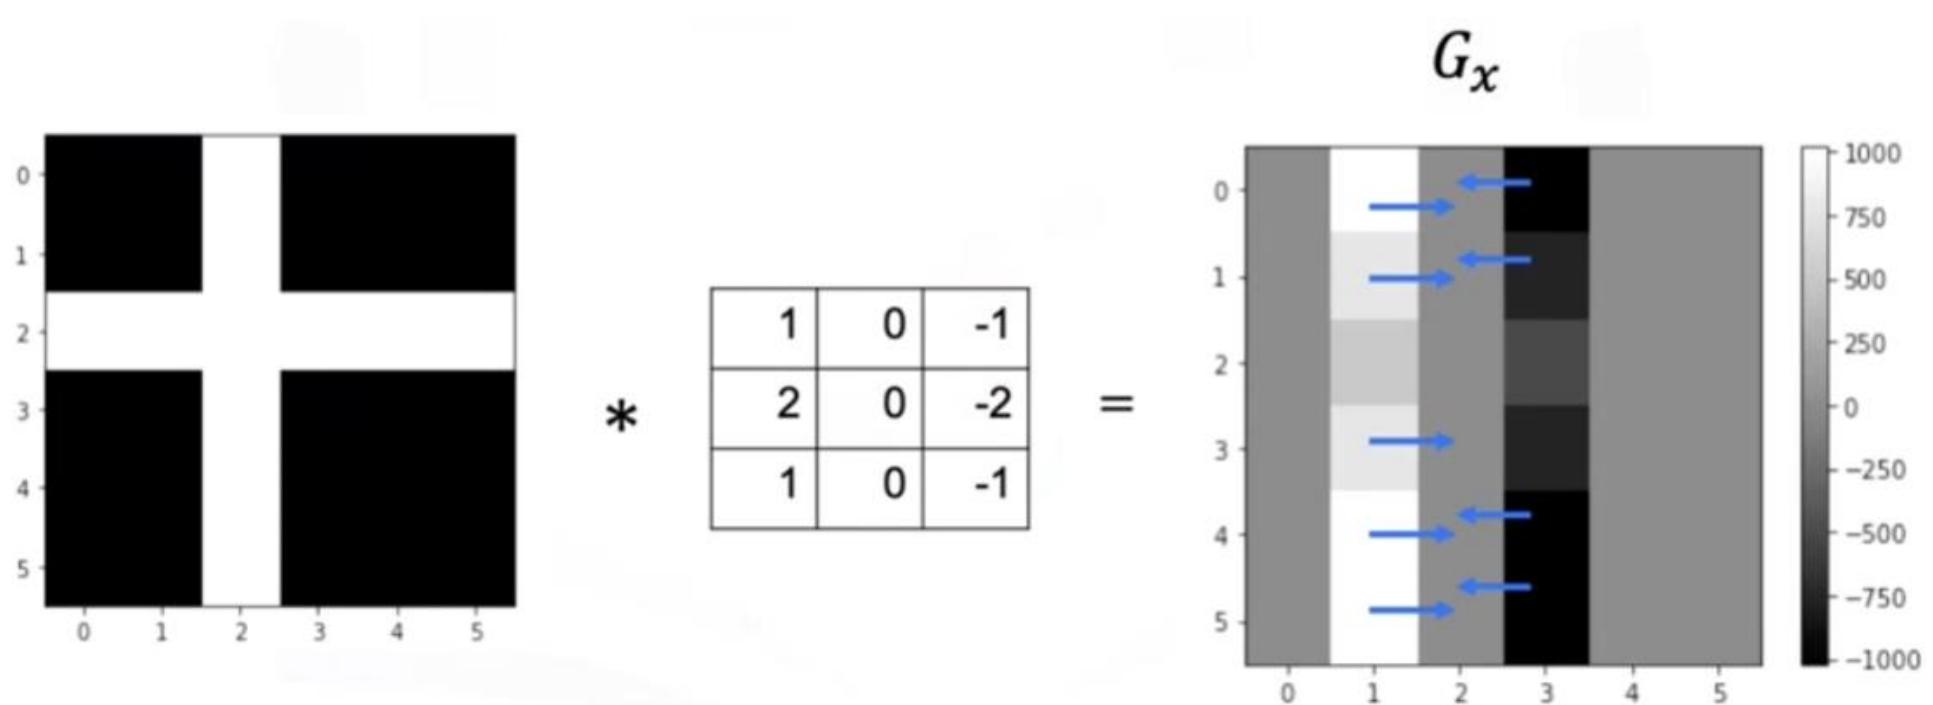
$$I[i, j + 1] - I[i, j]$$

$$\frac{\partial I}{\partial x}$$

i	$I[i, 2]$	$I[i, 1]$	$I[i, j + 1] - I[i, j]$
0	255	0	255
1	255	0	255
2	255	255	0
3	255	0	255
4	255	0	255
6	255	0	255



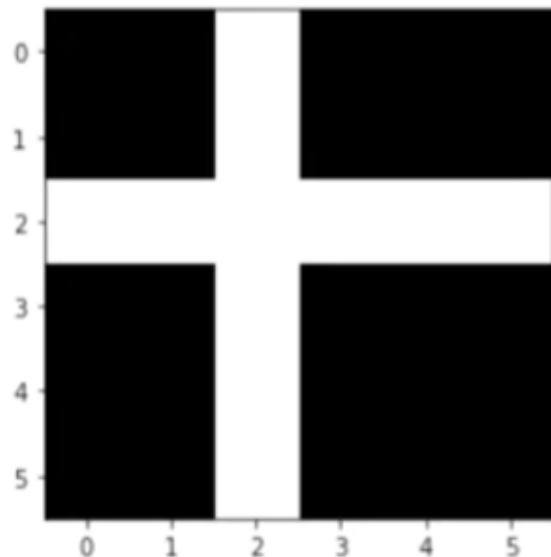
Edge Detection Filters



Horizontal Edge Detection



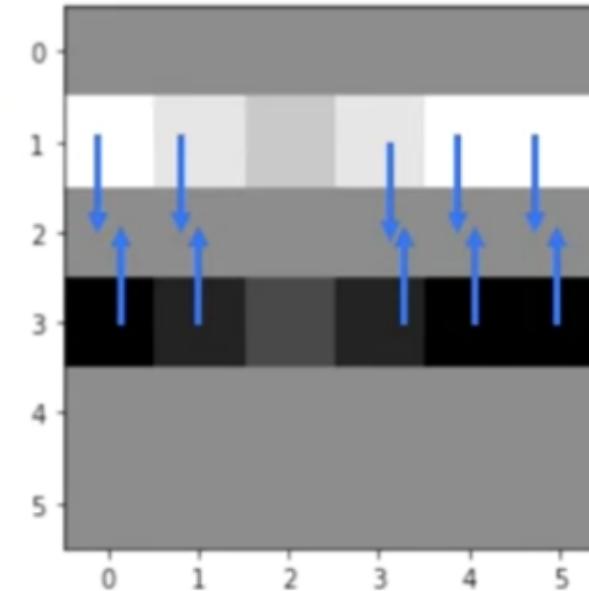
Edge Detection Filters



*

$$\begin{array}{|c|c|c|} \hline 1 & 2 & -1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

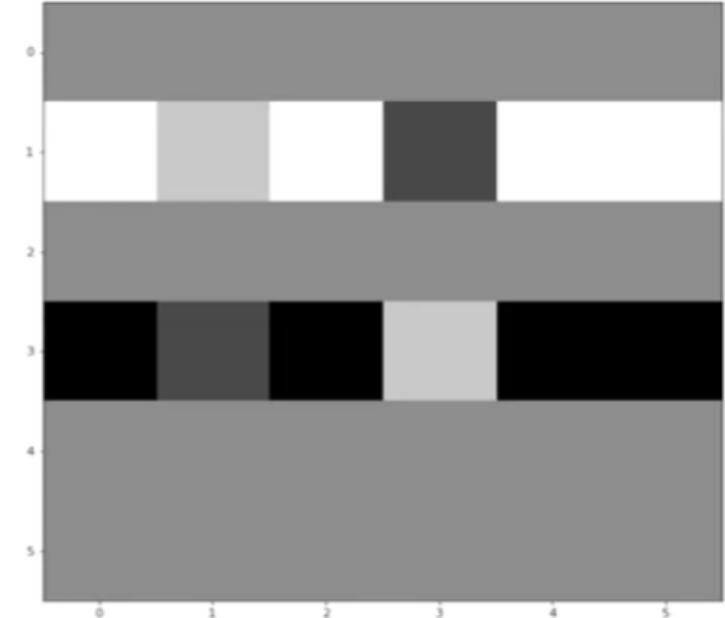
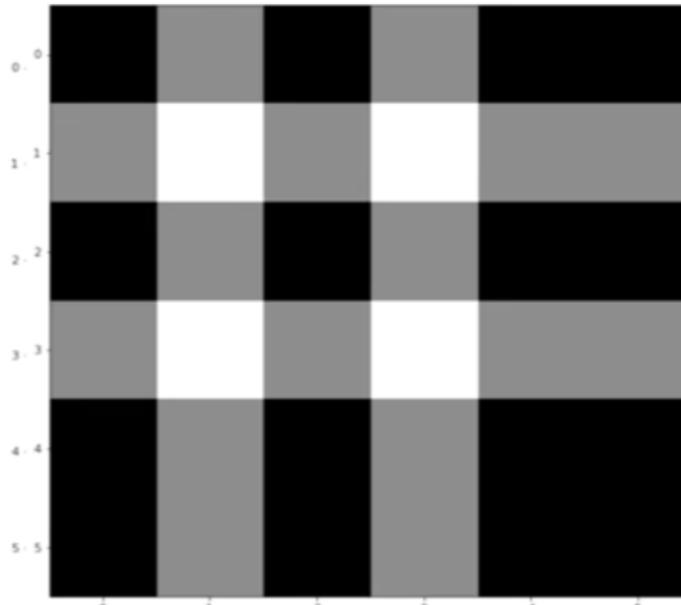
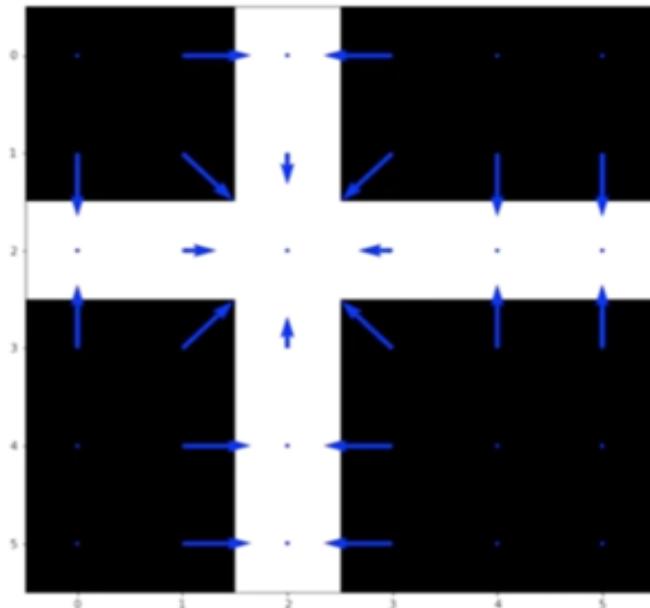
=



Vertical Edge Detection



Edge Detection Filters



$$\|G\|$$

$$G = [G_x, G_y]$$

$$\|G\| = \sqrt{{G_x}^2 + {G_y}^2}$$

$$\theta = \arctan\left(\frac{G_y}{G_x}\right)$$

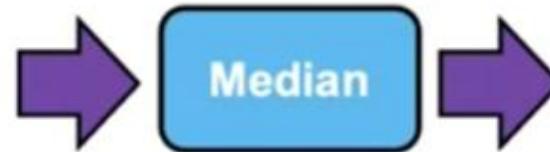


Median Filters – Non-Linear Filtering



Median filters are best for Impulse or Salt-and-Pepper noise.

255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
0	0	0	0	0	0
0	0	0	0	0	255



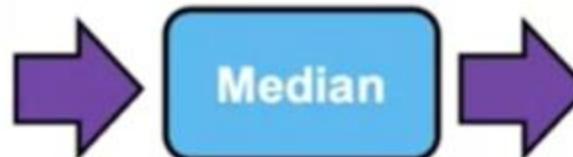
			255		



Median Filters – Non-Linear Filtering

Median filters are best for Impulse or Salt-and-Pepper noise.

255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
0	0	0	0	0	0
0	0	0	0	0	255



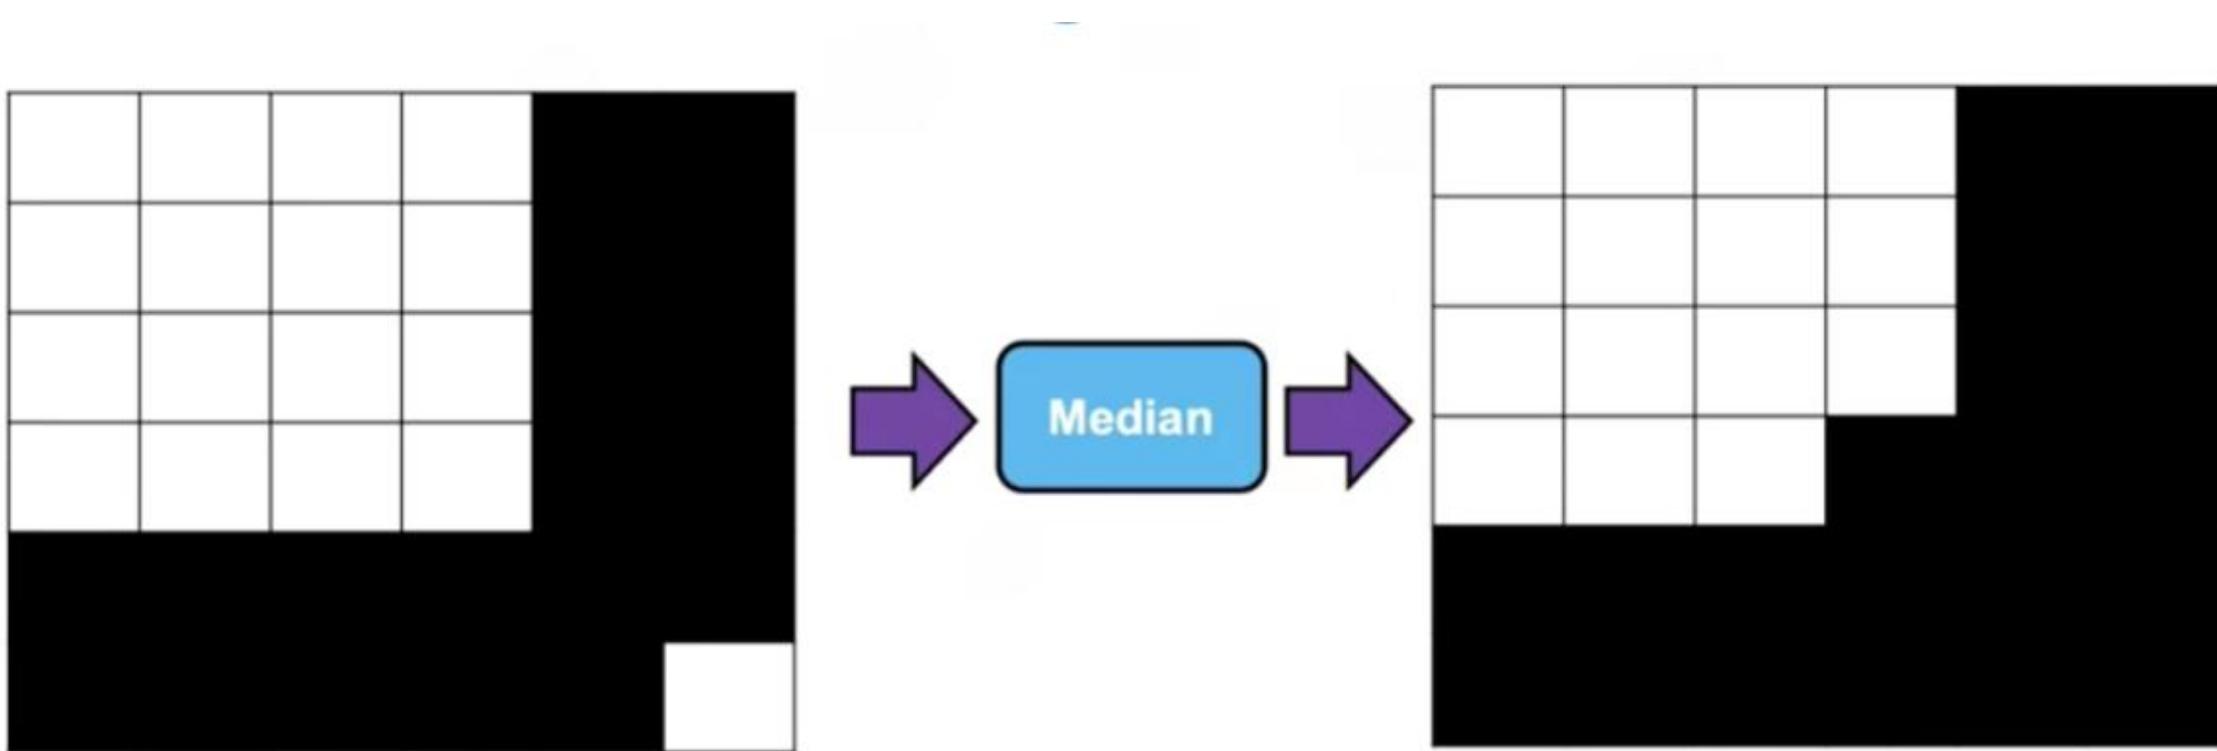
255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
255	255	255	255	0	0
0	0	0	0	0	0
0	0	0	0	0	0



Median Filters – Non-Linear Filtering



Median filters are best for Impulse or Salt-and-Pepper noise.
However, these may distort the shape.



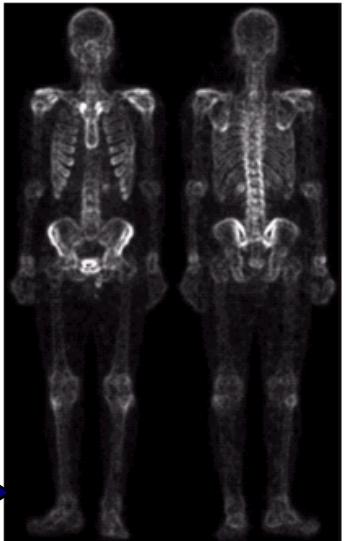


Using Multiple Filters



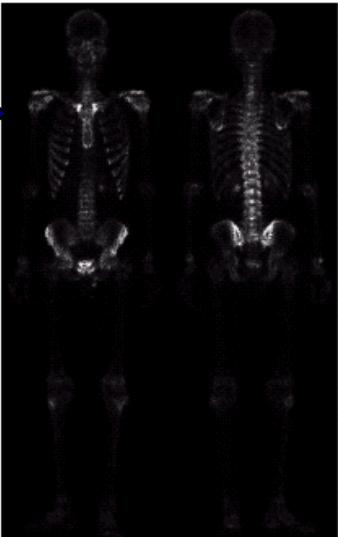
The product of (c)
and (e) which will be
used as a mask

(e)



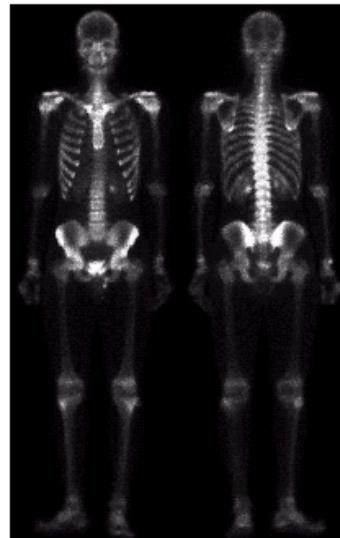
Sharpened image
which is sum of (a)
and (f)

(f)



Result of applying a
power-law trans. to
(g)

(g)



(h)

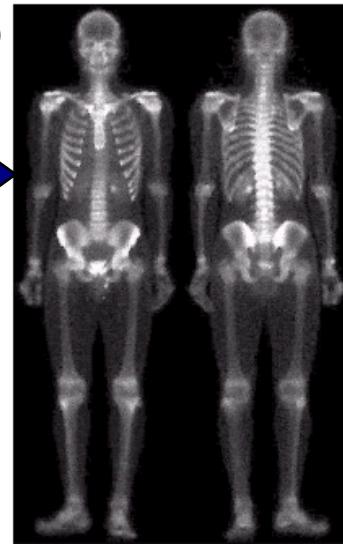


Image (d) smoothed with a 5×5 averaging filter

Combining Spatial Enhancement Methods



Acknowledgements



Some of the content of these slides is taken from:

- www.coursera.com
- Digital Image Processing by Gonzalez