

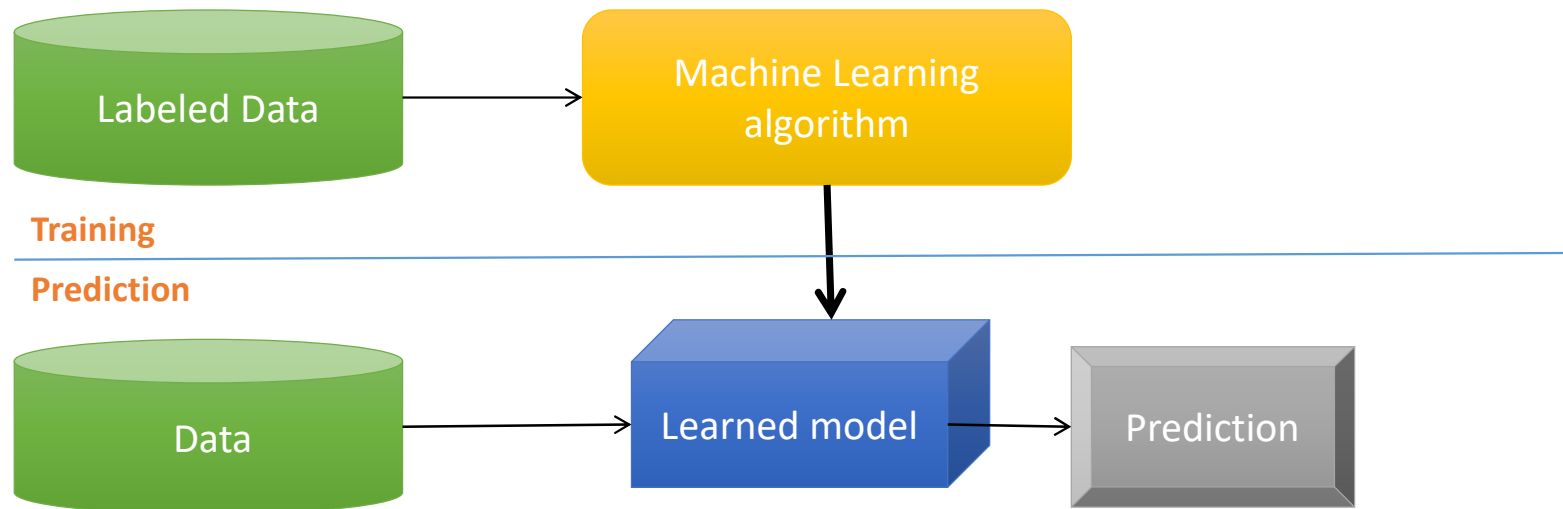
Neural Networks

Dr. Jameel Malik

muhammad.jameel@seecs.edu.pk

Machine Learning Basics -- Recap

Machine learning is a field of computer science that gives computers the ability to **learn without being explicitly programmed**




Methods that can learn from and make predictions on data

Why Neural Networks ?

- Let's say we need to learn complex non-linear hypothesis


Classification Problem



Why Neural Networks ?

- Suppose we need to learn complex non-linear hypothesis

Classification Problem




Assume the following Logistic Regression Hypothesis

$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1 x_2 + \theta_4 x_1^2 x_2 + \theta_5 x_1^3 x_2 + \theta_6 x_1 x_2^2 + \dots)$$

Why Neural Networks ?

- Suppose we need to learn complex non-linear hypothesis

Classification Problem



x_1 = size
 x_2 = # bedrooms
 x_3 = # floors
 x_4 = age
...
 x_{100}

Assume the following Logistic Regression Hypothesis


$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1 x_2 + \theta_4 x_1^2 x_2 + \theta_5 x_1^3 x_2 + \theta_6 x_1 x_2^2 + \dots)$$

- What should be the hypothesis when we have lots of features?

Why Neural Networks ?

- Suppose we need to learn complex non-linear hypothesis

Classification Problem



x_1 = size
 x_2 = # bedrooms
 x_3 = # floors
 x_4 = age
...
 x_{100}

Assume the following Logistic Regression Hypothesis

$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1 x_2 + \theta_4 x_1^2 x_2 + \theta_5 x_1^3 x_2 + \theta_6 x_1 x_2^2 + \dots)$$

- What should be the hypothesis when we have lots of features?
 - maybe we should assume a non-linear hypothesis with 1000's of features..
 - This can really blow up the features space
 - And we do not know which features can be more useful

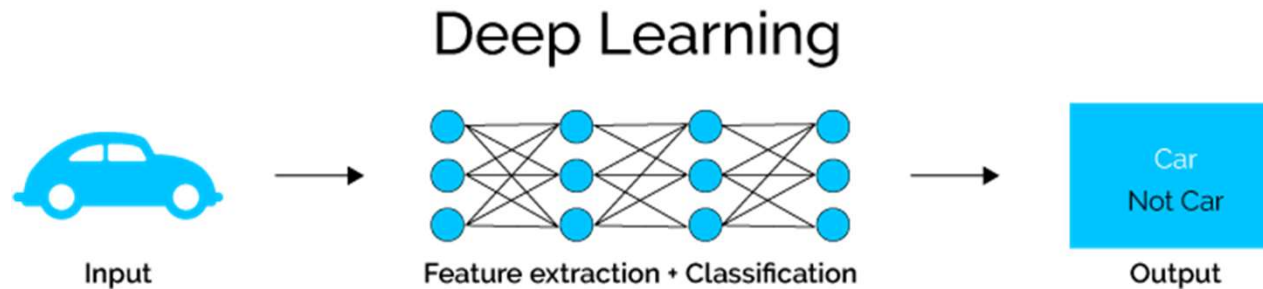

Introduction to Neural Networks

- Neural networks (NNs) is a type of machine learning algorithm that became popular in the 1980s.
 - Lots of successes, hype, and great conferences: NeurIPS, ICML etc.
 - Then along came SVMs, Random Forests and Boosting in the 1990s, and Neural Networks took a back seat.
 - Re-emerged around 2010 as **Deep Learning**.
 - By 2020s very dominant and successful.
 - Part of success due to vast improvements in computing power, larger training sets, and software: Tensorflow and PyTorch.
- Much of the credit goes to three pioneers and their students: Yann LeCun, Geoffrey Hinton and Yoshua Bengio, who received the 2019 ACM Turing Award for their work in Neural Networks.



What is Deep Learning (DL) ?

Deep learning algorithms attempt to learn (multiple levels of) representation by using a **hierarchy of multiple layers**



<https://www.xenonstack.com/blog/static/public/uploads/media/machine-learning-vs-deep-learning.png>


Why is DL useful?


- Manually designed features are often **over-specified**, **incomplete** and take a **long time to design** and validate
- Automatically learned Features are **easy to adapt**, **fast** to learn
- Deep learning provides a very **flexible** learnable framework for representing world, visual and linguistic information
- Can learn from both unsupervised and supervised data
- Effective **end-to-end** system learning
- Utilize large amounts of training data

In ~2010 DL started outperforming other ML techniques first in speech and computer vision, then Natural Language Processing (NLP).

Representational Power of NNs


- NNs with at least one hidden layer are universal approximators
i.e., NN can approximate any arbitrary complex continuous function
- NNs use **nonlinear mapping** of the inputs x to the outputs $f(x)$ to compute complex decision boundaries
- But then, why use **deeper** NNs?
 - The fact that deep NNs work better is an empirical observation






Perceptron



- The *perceptron* is the basic processing element of NN.
- It has inputs that may come from a sensor or may be the outputs of other perceptrons.




Perceptron

- The *perceptron* is the basic processing element of NN.
- It has inputs that may come from a sensor or may be the outputs of other perceptrons.





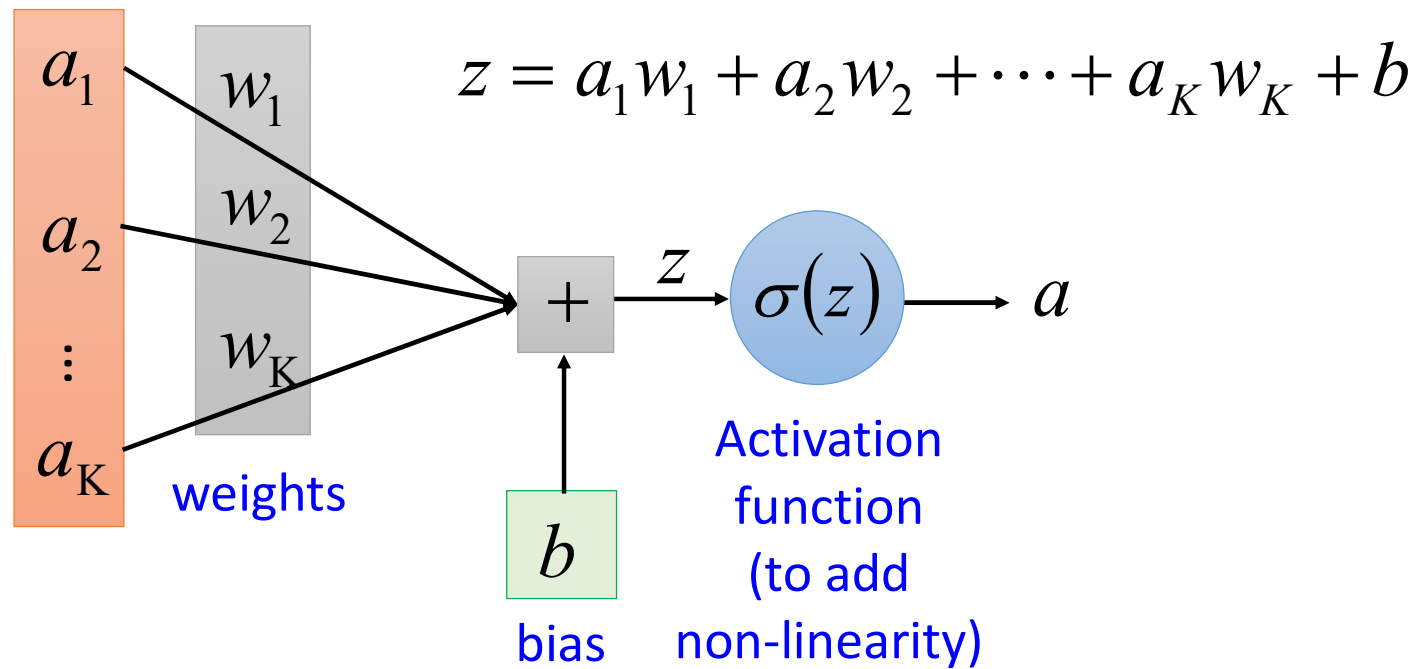
Nonlinear **sigmoid activation** function (f)



Elements of Neural Network


A single **neuron** maps a set of inputs into an output number, or $f: R^K \rightarrow R$

Neuron $f: R^K \rightarrow R$



A Simple Linear Layer

- **Memory Requirements:**
 - Parameters (weights, biases)
 - **12 + 3 = 15**









4 Input Neurons, 3 output Neurons

Thread 1

Thread 2


Thread 3



TIME

A Simple Linear Layer – CPU execution?

- **Memory Requirements:**
 - Parameters (weights, biases)
 - **12 + 3 = 15**









4 Input Neurons, 3 output Neurons

Thread 1

Thread 2


Thread 3



TIME

A Simple Linear Layer – GPU execution?

- **Memory Requirements:**
 - Parameters (weights, biases)
 - **12 + 3 = 15**









4 Input Neurons, 3 output Neurons

Thread 1


Thread 2

Thread 3




TIME

Example of Neural Network -- Prediction

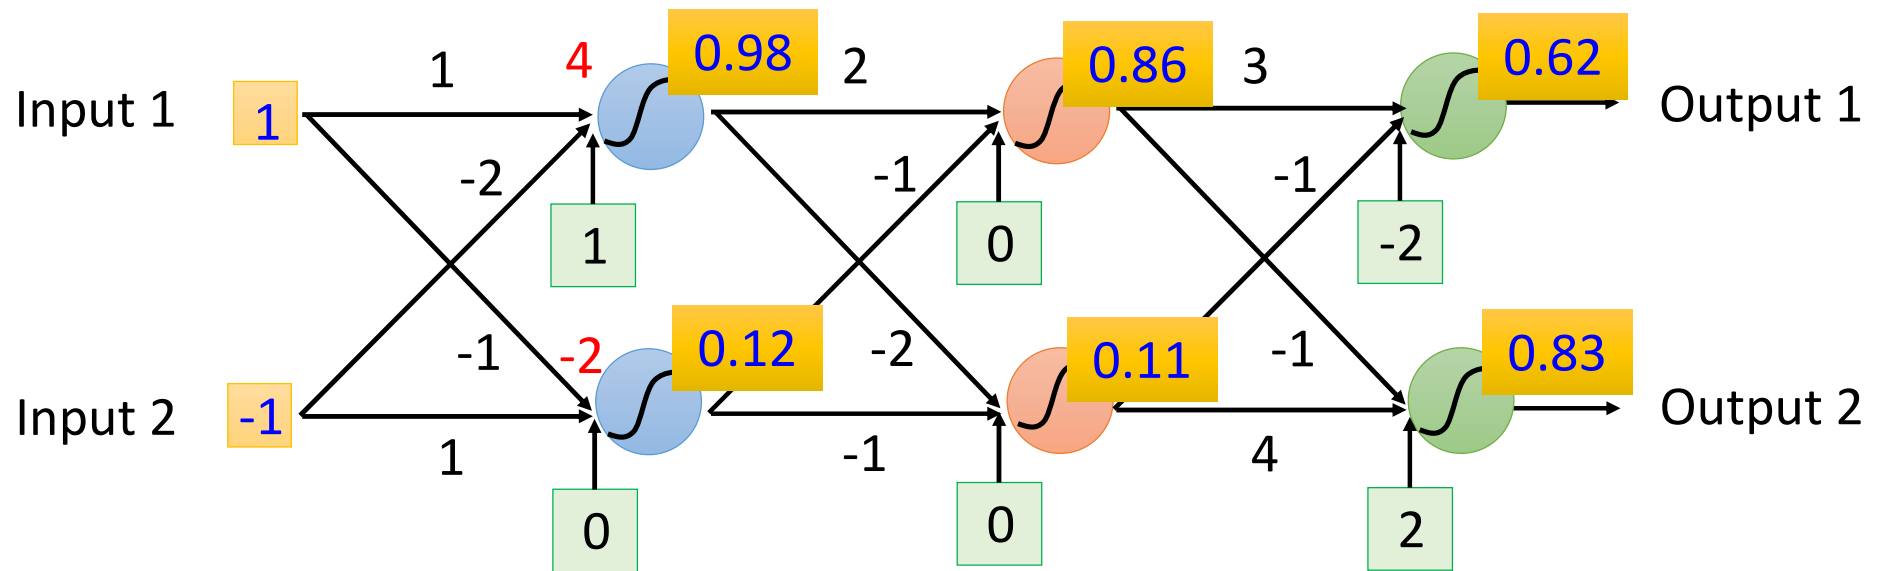


Sigmoid Function

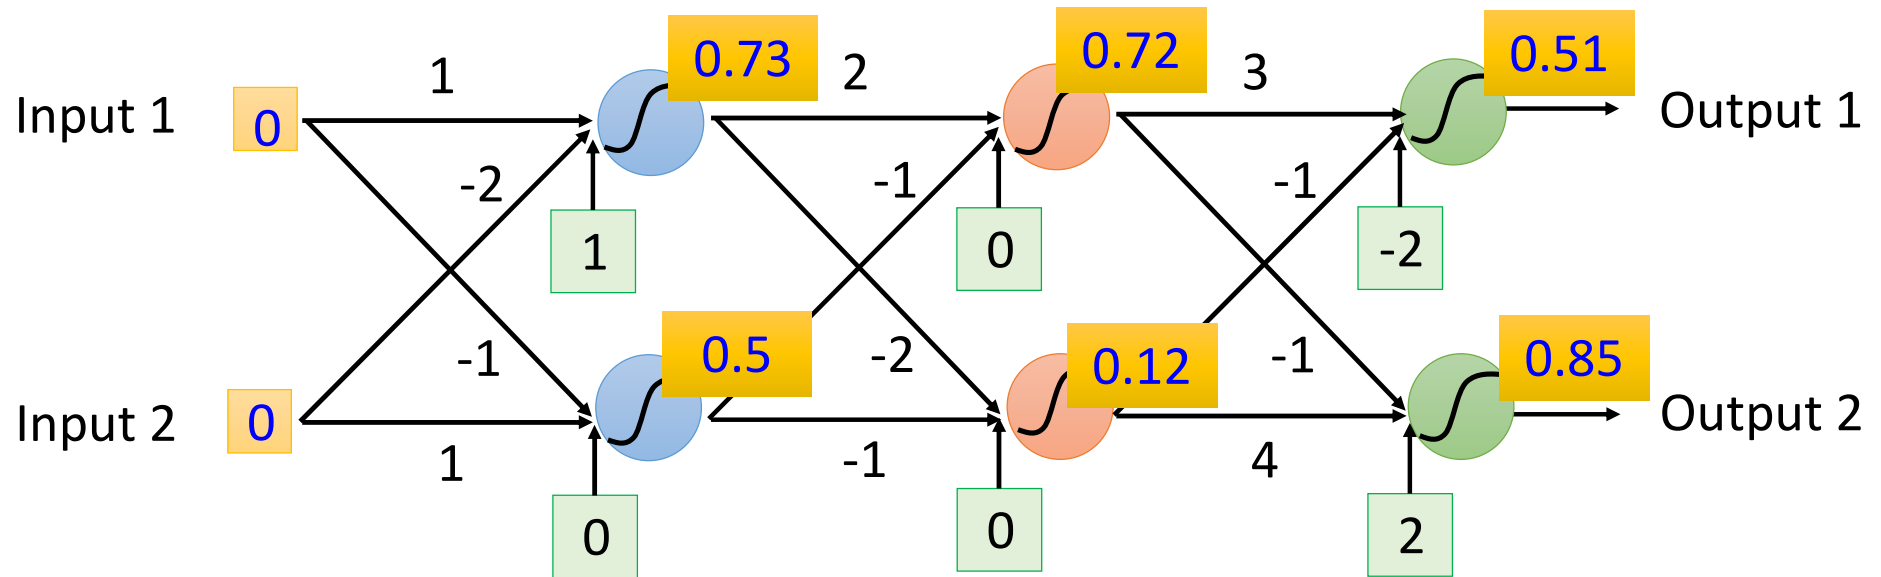
$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



Example of Neural Network -- Prediction




Example of Neural Network -- Prediction



$$f: \mathbb{R}^2 \rightarrow \mathbb{R}^2 \quad f\left(\begin{bmatrix} 1 \\ -1 \end{bmatrix}\right) = \begin{bmatrix} 0.62 \\ 0.83 \end{bmatrix} \quad f\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}\right) = \begin{bmatrix} 0.51 \\ 0.85 \end{bmatrix}$$

Different parameters define different function


Matrix Operation



$$\sigma\left(\underbrace{\begin{bmatrix} 1 & -2 \\ -1 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ -1 \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \end{bmatrix}}_{\begin{bmatrix} 4 \\ -2 \end{bmatrix}} \right) = \begin{bmatrix} 0.98 \\ 0.12 \end{bmatrix}$$

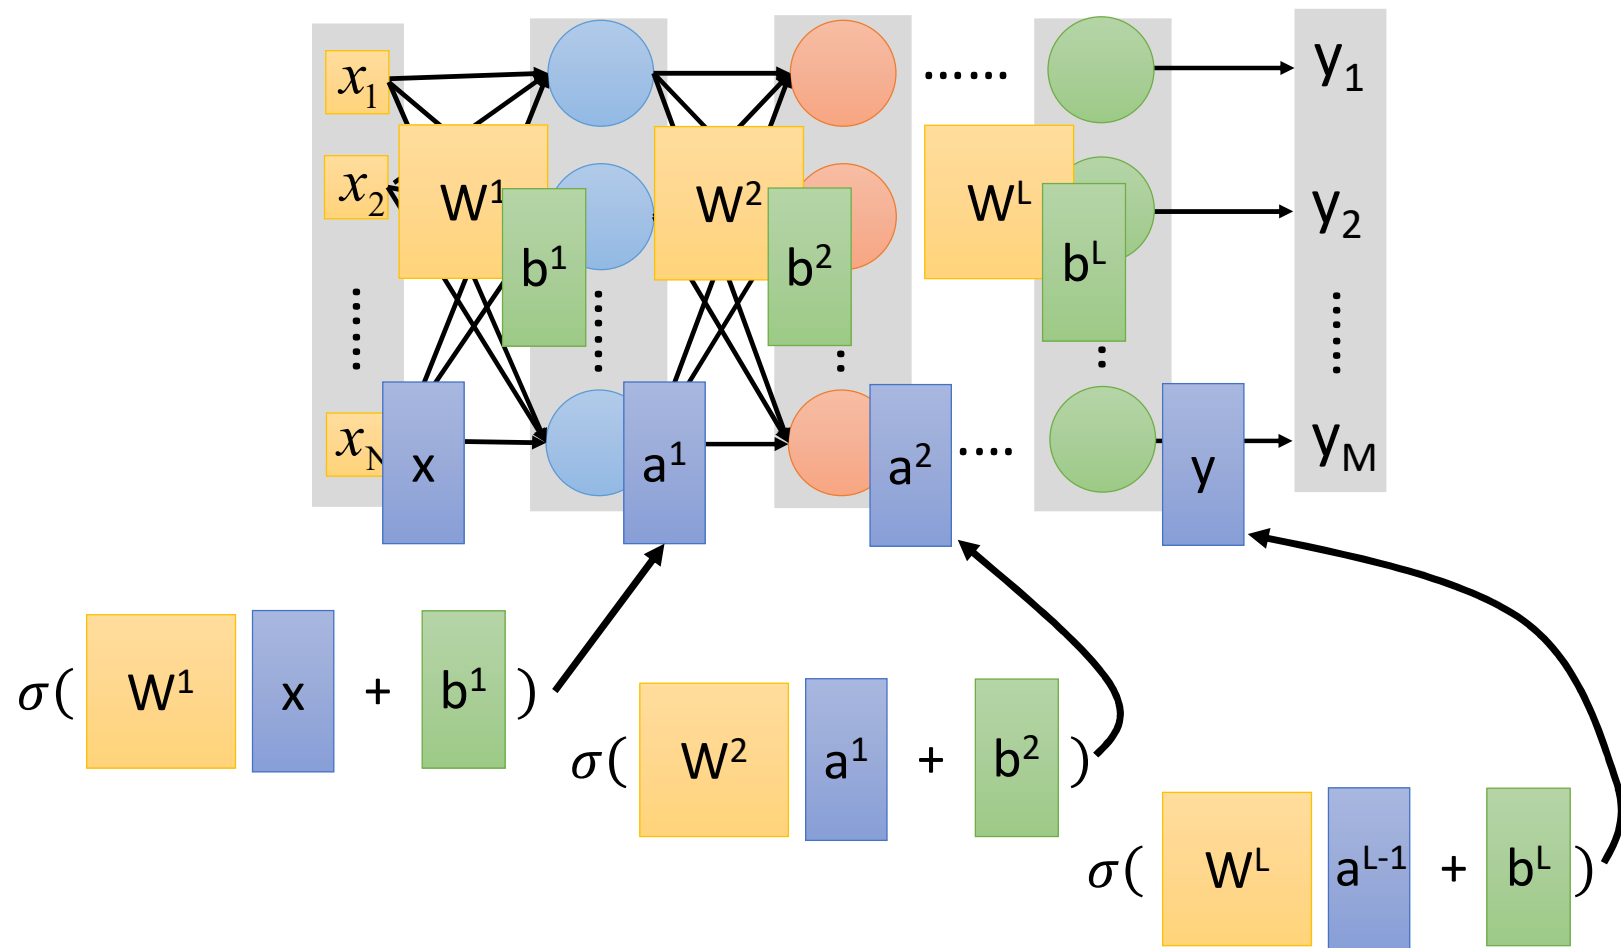
$$\sigma\left(\begin{bmatrix} W \end{bmatrix} \begin{bmatrix} x \end{bmatrix} + \begin{bmatrix} b \end{bmatrix} \right) = \begin{bmatrix} a \end{bmatrix}$$

Neural Network




Deep means many hidden layers

Neural Network



Neural Network



$$y = f(x)$$

Using parallel computing techniques
to speed up matrix operation

$$= \sigma(W^L \dots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \dots + b^L)$$