


Neural Networks


Dr. Jameel Malik

muhammad.jameel@seecs.edu.pk

Images to Digits - A Mapping from 784 to 10 Dimensions




**28x28
pixel
input
images
(784 dim)**


$$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ \vdots \end{bmatrix}$$

**output
vector
(10 dim)**

[Partial image courtesy: Nielsen]

Loss Function $L_i(\boldsymbol{\theta}) \mapsto \mathbb{R}$



Compare output layer to the true label

$$L_i(\boldsymbol{\theta}) = ||\text{output}_i - \text{label}_i||^2$$

The Parameters We Want


- Parameter θ^* that minimize the sum of avg. squared losses over all examples

$$\theta^* = \arg \min_{\theta} L(\theta) = \arg \min_{\theta} \sum_i ||f(\theta, x_i) - y_i||^2$$

- The squared loss is only one possible loss, several other options available
- **Goal:** Find the parameter vector θ^* for the labeled training set $\{(x_i, y_i)\}_{i=1}^I$ given the loss L

Gradient Descent Algorithm -- Recap

- Steps in the *gradient descent algorithm*:
 1. Randomly initialize the model parameters, θ^0
 2. Compute the gradient of the loss function at the initial parameters θ^0 : $\nabla\mathcal{L}(\theta^0)$
 3. Update the parameters as: $\theta^{new} = \theta^0 - \alpha\nabla\mathcal{L}(\theta^0)$
 - Where α is the learning rate
 4. Go to step 2 and repeat (until a terminating criterion is reached)




Convolutional Neural Networks (CNNs)

- *Convolutional neural networks* (CNNs) were primarily designed for image data
 - Capture the spatial/contextual information in images.
- CNNs have **less parameters** than MLPs (fully-connected layers)
- Example:
 - MLP sees or processes a flattened 1D vector (784 dimensions) of an MNIST image.
 - CNN sees or processes the original image of MNIST in 2D

Convolutional Neural Networks (CNNs)

- A convolutional filter slides (i.e., convolves) across the image

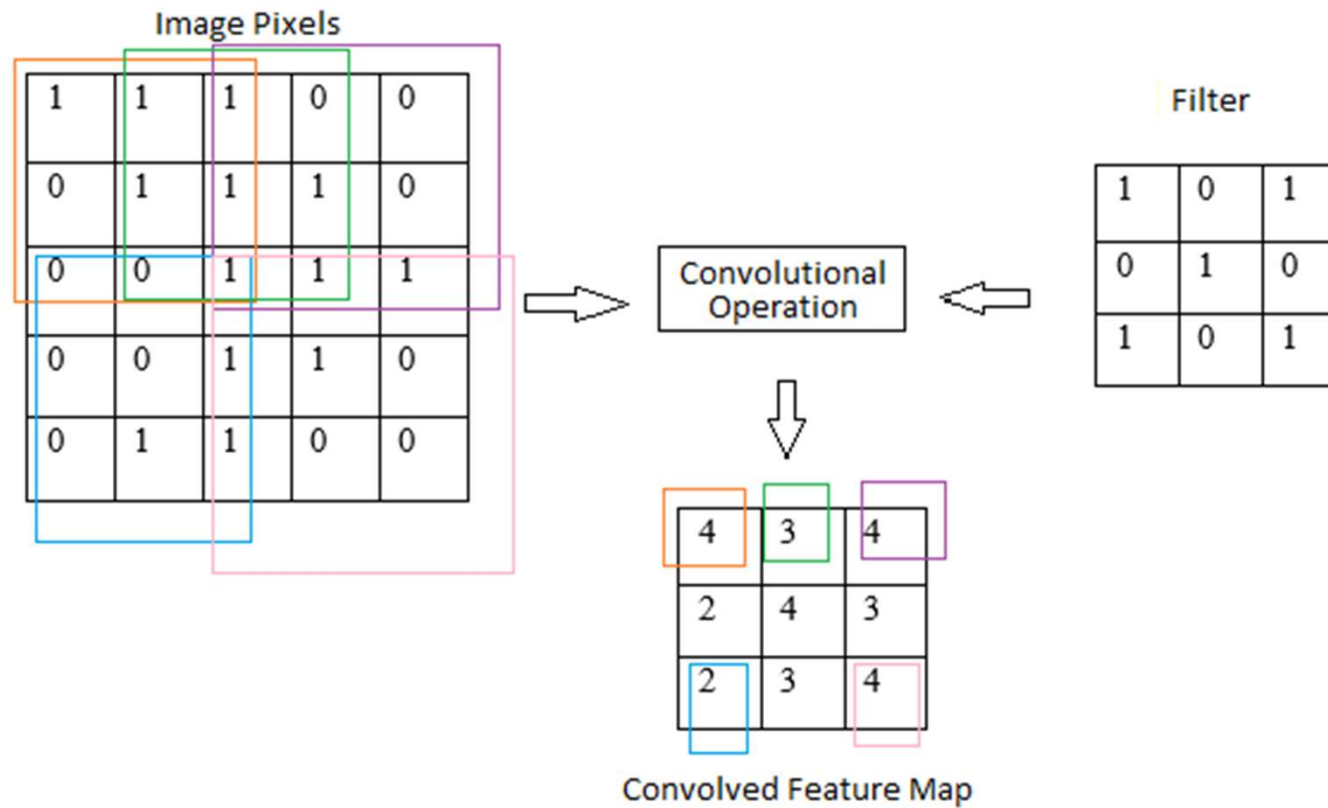


2D convolution

Picture from: http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution

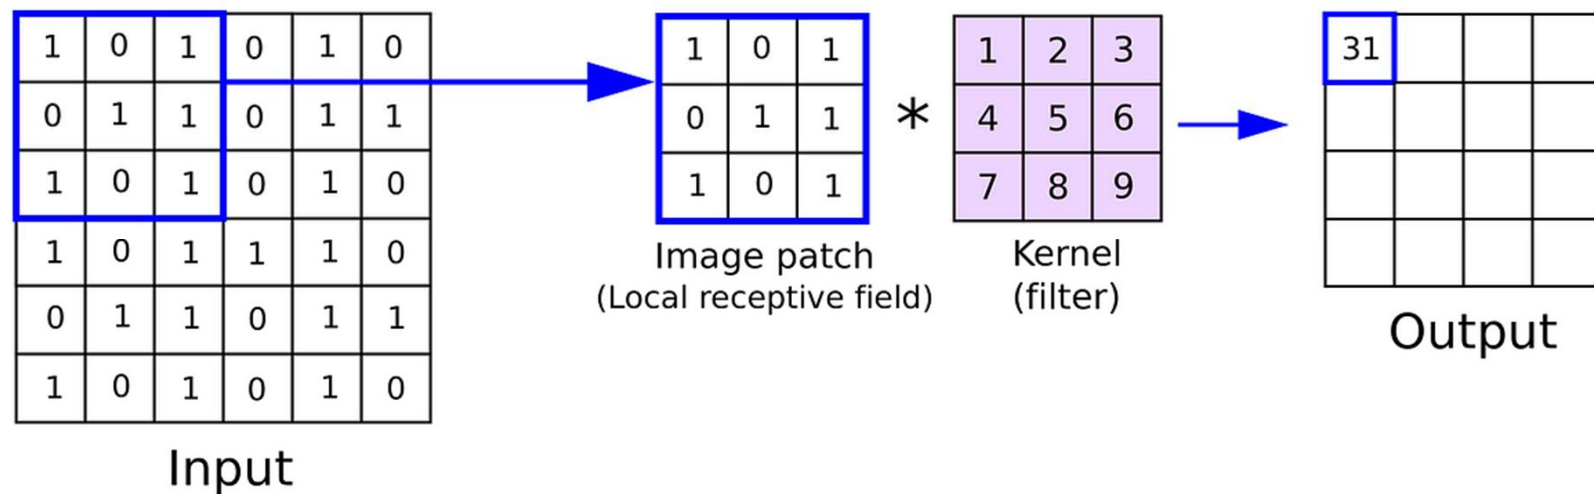
Convolutional Neural Networks (CNNs)

- A convolutional filter slides (i.e., convolves) across the image




Convolutional Neural Networks (CNNs)

- A convolutional filter slides (i.e., convolves) across the image



2D Convolution

$$y[m,n] = x[m,n] \otimes h[m,n] = \sum_{j=-\infty}^{\infty} \sum_{i=-\infty}^{\infty} x[i,j] \cdot h[m-i,n-j]$$



Laplacian Filter (3x3)

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

\otimes

Original Image



Laplacian filtered output



Convolution result
(activation/feature map)

Convolutional Neural Networks (CNNs)

Convolution Operation:

- When the convolutional filters are scanned over the image, they capture useful features
 - E.g., edge detection by convolutions

Filter 
$$\begin{pmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{pmatrix}$$



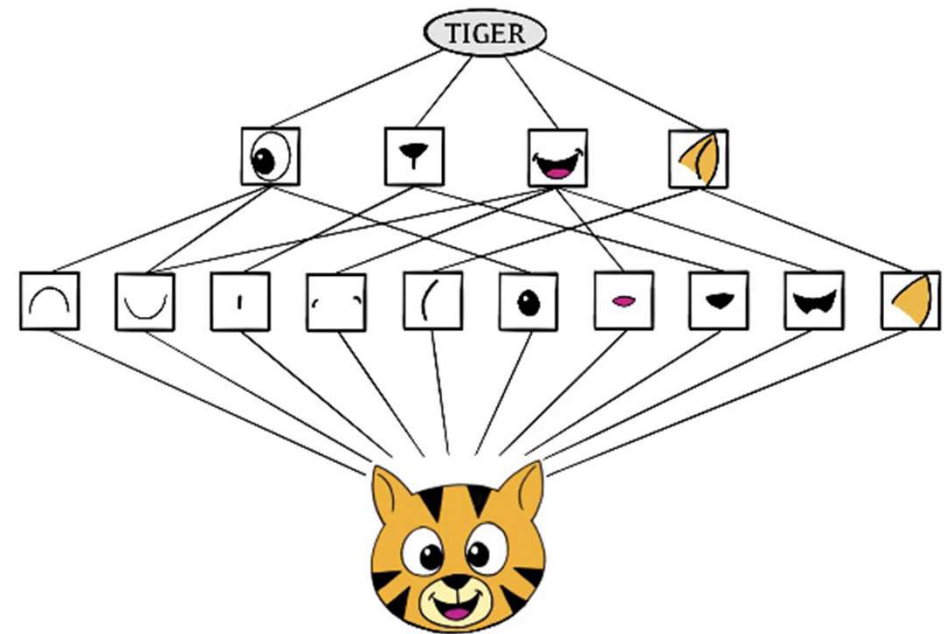
Input Image



Convolved
Image


How CNNs Work

- The CNN builds up an image in a hierarchical fashion.
- Edges and shapes (local features) are recognized and pieced together to form more complex shapes (compound features) such as eye and ear, eventually assembling the target image.
- This hierarchical construction is achieved using *convolution* layers.




Convolution Layer

32x32x3 image -> preserve spatial structure



Convolution Layer

- 32x32x3 image




- 5x5x3 filter



- **Convolve** the filter with the image
- i.e. “slide over the image spatially, computing dot products”

Convolution Layer

- 32x32x3 image



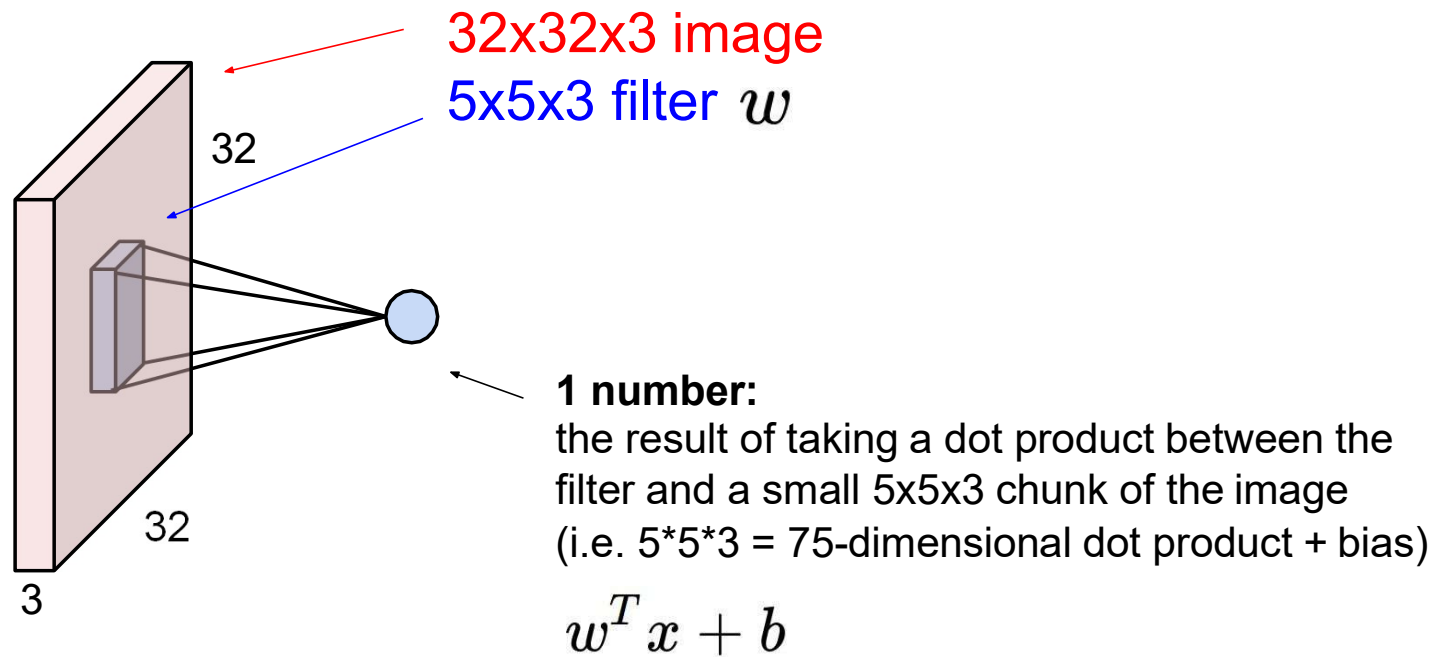
Filters always extend the full depth of the input volume

- 5x5x3 filter




- **Convolve** the filter with the image
- i.e. “slide over the image spatially, computing dot products”

Convolution Layer




Convolution Layer




1 number:
the result of taking a dot product between the
filter and a small 5x5x3 chunk of the image
(i.e. $5*5*3 = 75$ -dimensional dot product + bias)

$$w^T x + b$$





It's just a neuron with local connectivity...

Convolution Layer



The brain/neuron view of CONV Layer




An activation map is a 28x28 sheet of neuron outputs:

1. Each is connected to a small region in the input
2. All of them share parameters


“5x5 filter” -> “5x5 receptive field for each neuron”

Convolution Layer

consider a second, **green** filter




For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:




We stack these up to get a “new image” of size 28x28x6!




Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions



Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions




Animated AI



animatedai.github.io

two more layers to go: POOL/FC



Convolutional Neural Networks (CNNs)

Pooling layer:

- *Max pooling*: reports the maximum output within a rectangular neighborhood
- *Average pooling*: reports the average output of a rectangular neighborhood
- Pooling layers reduce the spatial size of the feature maps
 - Reduce the number of parameters, prevent overfitting

1	3	5	3
4	2	3	1
3	1	1	3
0	1	0	4

Input
Matrix


MaxPool with a 2×2 filter with stride of 2

4	5
3	4

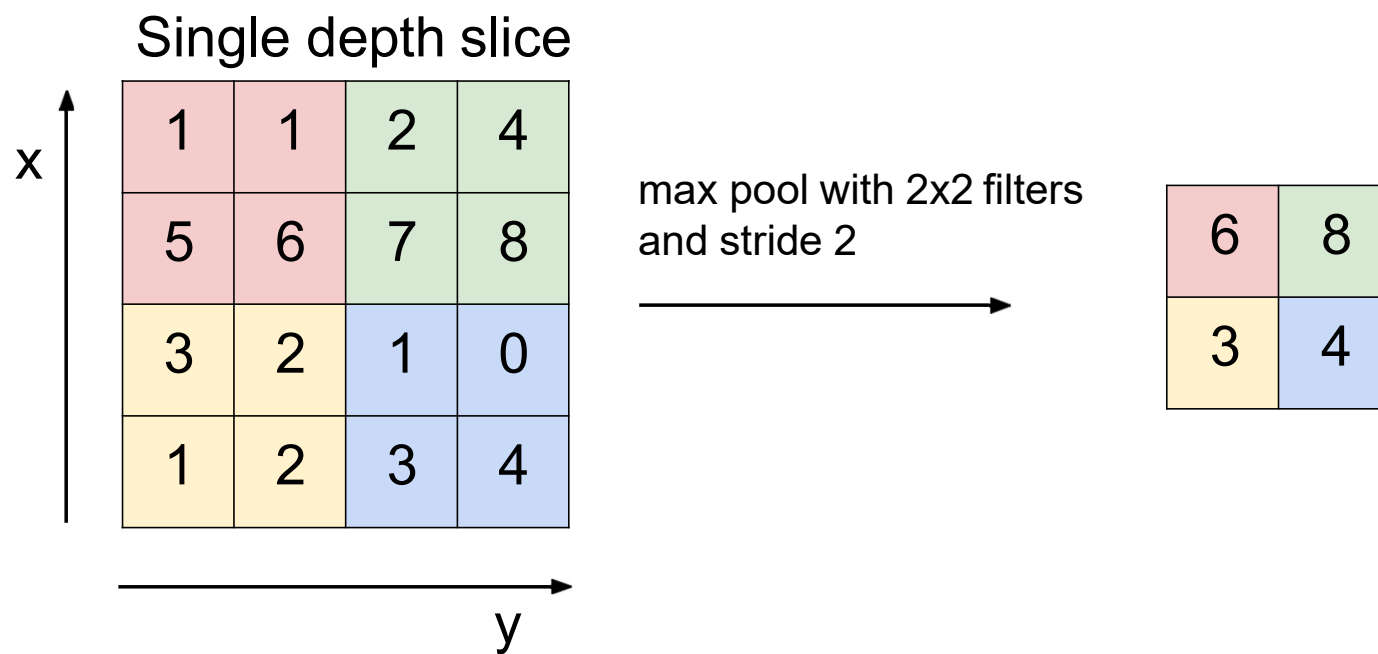
Output Matrix

Pooling layer


- makes the representations smaller and more manageable
- operates over each activation map independently:



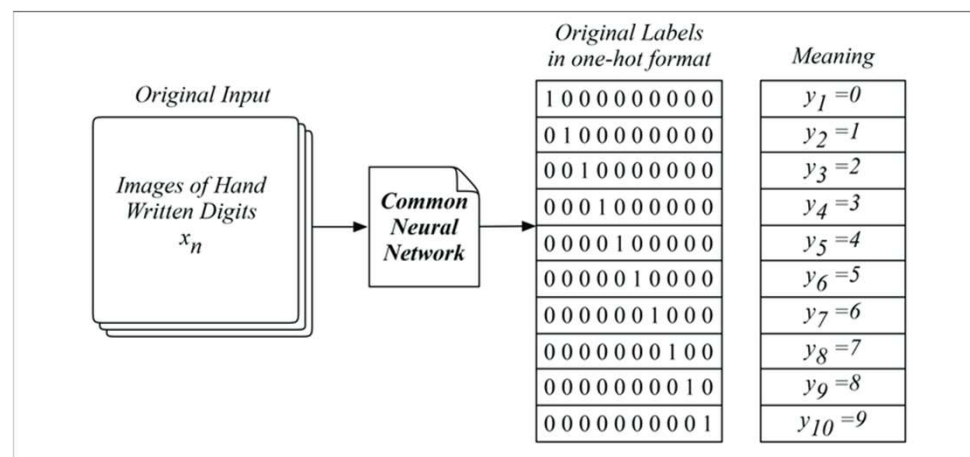
MAX POOLING



General CNN Architecture for MNIST Digit Classification

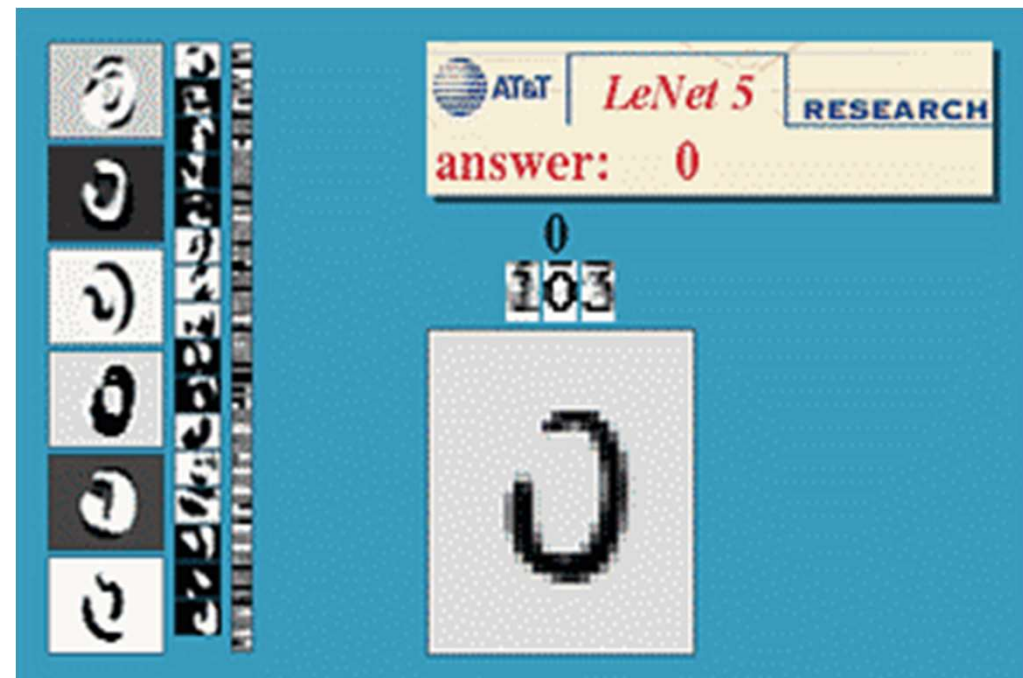



Inference



Training

General CNN Architecture for MNIST Digit Classification



General CNN Architecture for MNIST Digit Classification

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
conv2d_2 (Conv2D)	(None, 24, 24, 16)	4,624
conv2d_3 (Conv2D)	(None, 22, 22, 8)	1,160
max_pooling2d (MaxPooling2D)	(None, 11, 11, 8)	0
flatten (Flatten)	(None, 968)	0
dense (Dense)	(None, 10)	9,690

```
model = Sequential()
model.add(Input(shape=(28, 28, 1), name='input_layer'))
model.add(Conv2D(32, (3, 3), activation='relu', name='conv2d_1'))
model.add(Conv2D(16, (3, 3), activation='relu', name='conv2d_2'))
model.add(Conv2D(8, (3, 3), activation='relu', name='conv2d_3'))
model.add(MaxPooling2D(pool_size=(2, 2), name='max_pooling2d'))
model.add(Flatten(name='flatten'))
model.add(Dense(10, activation='softmax', name='dense'))
```